

PostgreSQL Storage Engine Architecture Atlas

21 Comprehensive Subsystem Chapters

Contents

1 PostgreSQL Storage Engine Architecture & Orientation

1. Executive Overview
2. Global Subsystem Invariants
3. Subsystem Dependency Matrix
4. PostgreSQL Fidelity Ledger

2 Item 1: Pager & 8KB Fixed Page Disk I/O

1. Architectural Role
2. Invariants & Formulas
3. Sequence Diagram: Reading & Writing Pages
4. PostgreSQL Fidelity Check

3 Item 2: Slotted Page Architecture

1. Physical Memory Layout
2. Invariants & Mathematical Formulas
3. Tuple Insertion Algorithm
4. PostgreSQL Fidelity Check

4 Item 3: Tuples & Heap CTID Physical Addressing

1. Physical CTID Tuple Addressing
2. Invariants & Binary Layout
3. Class Diagram: HeapFile & Pager Relationship
4. PostgreSQL Fidelity Check

5 Item 4: MVCC & Transaction Snapshot Visibility

1. Multi-Version Concurrency Control (MVCC) Overview
2. Invariants & Visibility State Machine
3. Sequence Diagram: Concurrent MVCC Reads and Writes
4. PostgreSQL Fidelity Check

6 Item 5: VACUUM Engine & In-Place Page Defragmentation

1. Dead Tuple Reclamation Mechanics
2. Invariants & CTID Stability
3. Sequence Diagram: Vacuum Page Lifecycle
4. PostgreSQL Fidelity Check

7 Item 6: Secondary B-Tree Index Subsystem

1. Index Decoupling & Multi-Version Addressing
2. Invariants & Lookup Pipeline
3. Sequence Diagram: Index Point Scan
4. PostgreSQL Fidelity Check

8 Item 7: Heap-Only Tuples (HOT) & Update Chains

1. The Write-Amplification Problem & HOT Solution

2. Invariants & Infomask Flags
3. Sequence Diagram: HOT Chain Traversal
4. PostgreSQL Fidelity Check

9 Item 8: Shared Buffers & Clock-Sweep Replacement

1. Buffer Pool Manager (Shared Buffers) Architecture
2. Invariants & Clock-Sweep State Machine
3. Sequence Diagram: Page Fetch & Eviction Lifecycle
4. PostgreSQL Fidelity Check

10 Item 9: Write-Ahead Logging (WAL) & REDO Durability

1. Write-Ahead Logging (WAL) Architecture
2. Invariants & Record Binary Layout
3. Sequence Diagram: REDO Crash Recovery Pass
4. PostgreSQL Fidelity Check

11 Item 10: TOAST Oversized-Attribute Storage

1. The TOAST Architecture
2. Invariants & Binary Layout
3. Sequence Diagram: Storing and Reassembling TOAST Data
4. PostgreSQL Fidelity Check

12 Item 11: SQL REPL CLI & End-to-End Engine Capstone

1. Engine Coordination Architecture
2. Invariants & Execution Rules
3. Sequence Diagram: End-to-End Insert Pipeline
4. PostgreSQL Fidelity Check

13 Item 12: Buffer Pool Single Gateway Invariant

1. The Single Gateway Principle
2. Invariants & Pin/Unpin Lifecycle State Machine
4. PostgreSQL Fidelity Check

14 Item 13: CLOG Commit Status 2-Bit Bitmap Pages

1. The Commit Log (CLOG) Architecture
2. Invariants & Mathematical Mapping Formulas
3. Sequence Diagram: CLOG Commit and Status Lookup
4. PostgreSQL Fidelity Check

15 Item 14: On-Disk B-Tree Index with Dynamic Node Splits

1. Disk-Resident B-Tree Architecture
2. Invariants & Binary Page Formats
3. Sequence Diagram: Leaf Split & Key Insertion
4. PostgreSQL Fidelity Check

16 Item 15: WAL Checkpoints & Full-Page Images (FPI)

1. Checkpoint & Torn-Page Protection
2. Invariants & Mechanics

3. Sequence Diagram: Checkpoint & Torn Page Healing
4. PostgreSQL Fidelity Check

17 Item 16: ARIES 3-Phase Crash Recovery & UNDO Pass (cpp-postgres only)

1. The 3-Phase ARIES Recovery Algorithm
2. Invariants & Compensation Log Records (CLRs)
3. Sequence Diagram: ARIES Crash Recovery Walkthrough
4. PostgreSQL Fidelity Check — why PostgreSQL has no UNDO

18 Item 17: TOAST Heap + Index Full Integration

1. Transparent Inline vs Out-of-Line Routing
2. Invariants & Infomask Flags
3. Sequence Diagram: Transparent TOAST Query Execution
4. PostgreSQL Fidelity Check

19 Item 18: Embedded HTTP/REST API Server (Approach A)

1. 2-Tier Architecture & Web Integration
2. Invariants & REST Endpoints
3. Sequence Diagram: React Browser Query via HTTP
4. PostgreSQL Fidelity Check

20 Item 19: PostgreSQL Wire Protocol Server (Approach C)

1. Enterprise 3-Tier Integration Architecture
2. Invariants & Packet Framing
3. Sequence Diagram: PostgreSQL Wire Protocol Handshake & Query
4. PostgreSQL Fidelity Check

21 Item 20: Native Desktop GUI & Unified Server Daemon

1. Unified Applications Overview
2. Desktop GUI Feature Breakdown
3. Sequence Diagram: Dual Server Daemon Operation
4. PostgreSQL Fidelity Check

1. PostgreSQL Storage Engine Architecture & Orientation

Confidence: verified

Citations: [include/pg/engine.h:1-85](#), [src/engine.cpp:1-250](#), [CMakeLists.txt:1-150](#)


1. Executive Overview

`cpp-postgres` is a zero-external-dependency relational storage engine implemented in C++20, written to teach PostgreSQL storage internals. It reproduces PostgreSQL's **mechanisms** — slotted pages, MVCC visibility rules, HOT chains, clock-sweep buffer replacement, WAL with full-page images, 2-bit CLOG bitmaps, and TOAST chunking — using deliberately narrowed struct layouts.

Implementation status. The behavioural audit in `notes/2026-08-27-architecture-audit.md` found that several mechanisms documented here existed as code but were not wired into the running system, and that write-ahead durability and the single-gateway rule were not upheld. Those are fixed as of 2026-08-27; each chapter carries an *Implementation status* section where its behaviour changed, and the audit holds the per-finding status table.

Fidelity disclaimer. This engine is *not* byte-compatible with a real PostgreSQL data directory, and several documented structures are smaller than PostgreSQL's. The single largest algorithmic divergence is crash recovery: `cpp-postgres` implements textbook **ARIES with an UNDO pass and CLRs**, whereas **PostgreSQL is redo-only and has no UNDO pass and no compensation log records at all** (see Item 16). Section 4 below is the authoritative fidelity ledger; every chapter ends with its own *PostgreSQL Fidelity Check*.

The architecture is partitioned into strict, decoupled layers:




2. Global Subsystem Invariants

The entire storage engine enforces seven architectural invariants:

- 8KB Page Uniformity:** All disk structures (Heap, B-Tree, CLOG, TOAST) operate on fixed 8,192-byte page boundaries aligned to the OS block layer (`[include/pg/constants.h:9]`). *PostgreSQL parity:* `BLCKSZ` defaults to 8192 but is a compile-time option (1 KB – 32 KB); `cpp-postgres` hard-codes 8192.
- Buffer Pool Single Gateway:** Application subsystems never execute direct file I/O; all page accesses pin frames inside `BufferPoolManager` (`[src/heap.cpp:45]`).
- Write-Ahead Logging (WAL):** A dirty page in RAM is never written to disk until its corresponding WAL log record has been flushed (`page.pd_lsn <= wal.flushed_lsn`) (`[src/wal.cpp:110]`).
- Append-Only MVCC:** Updates never rewrite a tuple's *payload* in place; they append a new version stamped with `xmin` and mark the old version by writing `xmax` into its header. Note that the header write **is** an in-place page mutation — which is precisely why the delete/update must be WAL-logged (`[src/heap.cpp:88]`).
- Secondary Index Decoupling:** B-Tree secondary indexes map `Key -> CTID` physical addresses without embedding tuple columns or MVCC headers (`[src/btree.cpp:15]`).
- Zero Index-Bloat HOT Updates:** Unindexed column updates residing on the same 8KB page construct line-pointer chains with zero index writes (`[src/heap.cpp:142]`).
- 2-Bit CLOG Density:** Transaction commit states are packed into 2-bit bitmaps on disk, scaling to 32,768 transactions per 8KB page (`[src/clog.cpp:30]`). *PostgreSQL parity:* exact — `CLOG_XACTS_PER_PAGE` is also 32,768, and the four status codes match bit-for-bit.

3. Subsystem Dependency Matrix



4. PostgreSQL Fidelity Ledger

Verified against PostgreSQL master headers (`bufpage.h` , `itemid.h` , `htup_details.h` , `heaptost.h` , `clog.h` , `xlogrecord.h` , `nbtree.h`).

4.1 Structures that match PostgreSQL exactly

Structure	Value	PostgreSQL source
Line pointer size	4 bytes	<code>ItemIdData</code>
Line-pointer state codes	<code>UNUSED=0, NORMAL=1, REDIRECT=2, DEAD=3</code>	<code>LP_UNUSED/LP_NORMAL/LP_REDIRECT/LP_DEAD</code>
<code>t_ctid</code> size	6 bytes	<code>ItemPointerData</code>
<code>PD_ALL_VISIBLE</code>	<code>0x0004</code>	<code>bufpage.h</code>
CLOG transactions per page	32,768	<code>CLOG_XACTS_PER_PAGE</code>
CLOG status codes	<code>00/01/10/11</code> = in-progress/committed/aborted/sub-committed	<code>clog.h</code>
<code>HEAP_HOT_UPDATED</code> / <code>HEAP_ONLY_TUPLE</code>	<code>0x4000</code> / <code>0x8000</code>	<code>htup_details.h</code> (but in <code>t_infomask2</code> , not <code>t_infomask</code>)
Max heap tuples per 8KB page	291	<code>MaxHeapTuplesPerPage</code> — same number, reached by different arithmetic
pgwire protocol	v3.0, <code>SSLRequest</code> code 80877103	<code>pqcomm.h</code>

4.2 Structures that are narrower than PostgreSQL's

Structure	cpp-postgres	PostgreSQL	Missing
Page header	18 B	24 B (<code>SizeOfPageHeaderData</code>)	<code>pd_page_size_version</code> (2 B), <code>pd_prune_xid</code> (4 B)
Line-pointer bit split	16 / 14 / 2 (offset / len / flags)	15 / 2 / 15 (<code>lp_off</code> / <code>lp_flags</code> / <code>lp_len</code>)	—
Tuple header	16 B	23 B (<code>SizeOfHeapTupleHeader</code> , MAXALIGNED to 24)	<code>t_cid</code> union (4 B), <code>t_infomask2</code> (2 B), <code>t_hoff</code> (1 B), <code>t_bits</code> null bitmap
Infomask	one 16-bit <code>infomask</code>	two words: <code>t_infomask</code> + <code>t_infomask2</code>	16 visibility/lock flags incl. all hint bits
<code>HEAP_HASEXTERNAL</code>	<code>0x2000</code>	<code>0x0004</code> (<code>0x2000</code> is <code>HEAP_UPDATED</code>)	—
WAL record header	35 B, custom	24 B <code>XLogRecord</code> + rmgr block headers	resource managers, block references
WAL checksum	CRC-32 (IEEE 802.3)	CRC-32C (Castagnoli)	—
TOAST threshold	2048 B per attribute	2032 B per whole tuple (<code>TOAST_TUPLE_THRESHOLD</code>)	—

Structure	cpp-postgres	PostgreSQL	Missing
TOAST chunk size	2048 B	1996 B (<code>TOAST_MAX_CHUNK_SIZE</code>)	—
Buffer pool	16 frames	16,384 frames (128 MB <code>shared_buffers</code> default)	freelist, strategy rings, usage-count cap of 5
B-tree node	parent pointer + right sibling	Lehman & Yao: metapage, high key, <code>btpo_prev</code> / <code>btpo_next</code> , no parent pointers	deduplication, fillfactor splits

4.3 PostgreSQL mechanisms this engine does not model at all

- **Transaction ID wraparound and freezing** (`FrozenTransactionId` , `vacuum_freeze_min_age` , anti-wraparound autovacuum). `tx_id_t` here is a plain monotonic `uint32_t` that is never frozen.
- **Hint bits** (`HEAP_XMIN_COMMITTED` , `HEAP_XMAX_COMMITTED` , ...) — PostgreSQL caches CLOG lookups in the tuple header; this engine hits CLOG on every visibility test.
- **Visibility map / free space map**, and therefore **index-only scans**.
- **MultiXactIds** and row-level locking (`SELECT ... FOR UPDATE` also writes `xmax` in PostgreSQL).
- **Sub-transactions / savepoints**, despite CLOG reserving the `SUB_COMMITTED` code.
- **Autovacuum, the checkpoint, the WAL writer, and the background writer** as separate processes.
- **Streaming replication, logical decoding, replication slots**.
- **Isolation levels**. PostgreSQL defaults to Read Committed (new snapshot per statement) and offers Repeatable Read and Serializable (SSI); this engine implements one snapshot-per-transaction model only.
- **The query planner and executor** — no cost model, no join methods, no statistics.


2. Item 1: Pager & 8KB Fixed Page Disk I/O

Confidence: verified

Citations: [include/pg/pager.h:1-62](#), [src/pager.cpp:1-85](#), [tests/test_pager.cpp:1-95](#)

1. Architectural Role


The **Pager** is the lowest-level storage primitive in the database engine. It abstracts the host operating system's filesystem into a linear sequence of fixed-size **8,192-byte (8KB)** disk pages indexed by a 32-bit `page_id_t`.



2. Invariants & Formulas

- Exact Page Boundary Alignment:** $\text{file_offset} = \text{page_id} \times 8192$ Every disk read and write transfers exactly 8,192 bytes. No partial page I/O is ever permitted.
- Deterministic File Growth:** When `write_page(page_id, ...)` is called on an unallocated `page_id \geq \text{num\_pages()}`, the file automatically grows in discrete 8KB increments padded with zeros (`[src/pager.cpp:52]`).
- C++ Stream Flag Integrity:** Because reading past EOF in `std::fstream` sets `eofbit` and `failbit`, subsequent seek operations silently fail unless `stream.clear()` is invoked prior to any seek or write (`[src/pager.cpp:25]`).

3. Sequence Diagram: Reading & Writing Pages



4. PostgreSQL Fidelity Check

PostgreSQL's counterpart to the Pager is the **storage manager** (`smgr`), whose only implementation is `md.c` ("magnetic disk").

Claim	Verdict
Relation file addressed as a linear array of fixed 8KB blocks	Exact
<code>offset = block_number * BLCKSZ</code> , whole blocks only, never partial I/O	Exact
Writing past the end extends the file in whole-block increments	Exact in shape — PostgreSQL's <code>mdextend()</code> zero-fills
<code>stream.clear()</code> before seek after EOF	C++-specific — an artifact of <code>std::fstream</code> , no PostgreSQL analogue
<code>flush()</code> after every write	Divergent — see below

What PostgreSQL layers on top

- **Segmentation.** A relation larger than 1 GB is split into `<relfilenode>` , `<relfilenode>.1` , `<relfilenode>.2` , ... Block 200,000 lives in segment 1 at an offset computed modulo `RELSEG_SIZE` . This engine uses one unbounded file.
- **Forks.** Each relation has a *main* fork plus `_fsm` (free space map), `_vm` (visibility map) and, for unlogged tables, `_init` . Only the main fork exists here.
- **fsync is deferred, not per-write.** PostgreSQL writes with buffered `pwrite()` and hands the fsync obligation to the checkpoint through a shared request queue (`register_dirty_segment`), so durability is paid once per checkpoint rather than once per page. Flushing on every `write_page()` , as here, is far more conservative — and far slower — than PostgreSQL.
- **relfilenode vs OID.** Filenames are `relfilenode` numbers, which *change* under `VACUUM FULL` , `TRUNCATE` and `REINDEX` ; they are not the table's OID.
- **Direct I/O and readahead.** `io_method` / `debug_io_direct` , and from v17 a streaming read interface that batches sequential block requests.

Implementation status (2026-08-27)

The Pager no longer uses `std::fstream` . It sits on a raw descriptor (`pg::File`) with positioned reads and writes, which gives it two things a stream cannot:

- **A real durability barrier.** `Pager::sync()` calls `fdatsync` (`_commit` on Windows). `std::fstream::flush()` only drains the userspace buffer into the OS page cache, so before this the WAL was not durable at all.
- **No sticky stream state.** Positioned I/O has no shared cursor, so a read that hits end-of-file cannot silently suppress a later write. That failure mode had been quietly discarding the WAL records written during crash recovery.

Still simplified: one file per relation with no 1GB segmentation, no forks, and `fsync` at checkpoint rather than through a checkpoint's request queue.


3. Item 2: Slotted Page Architecture

Confidence: verified

Citations: [include/pg/page.h:1-125](#), [src/page.cpp:1-170](#), [tests/test_page.cpp:1-110](#)

1. Physical Memory Layout

The **Slotted Page** architecture solves variable-length record management and in-place defragmentation. Within each 8,192-byte page, memory grows inward from both ends:



2. Invariants & Mathematical Formulas

1. Page Header Structure (18 Bytes — PostgreSQL uses 24) (`[include/pg/page.h:41-48]`):

- `pd_lsn` (8 Bytes): LSN of the last WAL record that modified this page (`[include/pg/page.h:42]`).
- `pd_checksum` (2 Bytes): Page checksum (`[include/pg/page.h:43]`).
- `pd_flags` (2 Bytes): Page status flags (e.g. `PD_ALL_VISIBLE = 0x0004`).
- `pd_lower` (2 Bytes): Byte offset marking the end of the line pointers array (initially 18).
- `pd_upper` (2 Bytes): Byte offset marking the start of the youngest tuple data (initially 8192).
- `pd_special` (2 Bytes): Reserved for index-specific metadata (8192 for heap pages — same convention as PostgreSQL, where a heap page has no special area).

PostgreSQL's `PageHeaderData` is **24 bytes** (`SizeOfPageHeaderData`). The first six fields are identical in order and width; PostgreSQL then adds two fields this engine omits:

- `pd_pagesize_version` (2 Bytes): page size in the high byte, layout version in the low byte — the check that refuses to mount a data directory built with a different `BLCKSZ` .
- `pd_prune_xid` (4 Bytes): hint holding the oldest XID on the page that might be prunable. This is what lets an ordinary `SELECT` decide cheaply whether HOT pruning is worth attempting (Item 7).

`pd_flags` values match PostgreSQL exactly: `PD_HAS_FREE_LINES = 0x0001` , `PD_PAGE_FULL = 0x0002` , `PD_ALL_VISIBLE = 0x0004` .

2. Line Pointer Layout (4 Bytes) (`[include/pg/page.h:22-38]`):

- `lp_offset` (`uint16_t`, 16 bits): Byte offset from start of page where tuple data begins.

- `lp_len_flags` (`uint16_t`, 16 bits):
 - Lower 14 bits (`0x3FFF`): Byte length of tuple.
 - Upper 2 bits (`>> 14`): `ItemFlags` (`00` = `UNUSED`, `01` = `NORMAL` (Live), `10` = `REDIRECT` (HOT chain), `11` = `DEAD`).

The **size (4 bytes) and the four state codes are identical to PostgreSQL**, but the bit split is not.

PostgreSQL's `ItemIdData` is a single 32-bit bitfield triple:

```
unsigned lp_off:15,    /* offset to tuple from start of page */
        lp_flags:2,   /* LP_UNUSED / LP_NORMAL / LP_REDIRECT / LP_DEAD */
        lp_len:15;    /* byte length of tuple */
```

15 bits is exactly enough to address any byte in a 32 KB page, which is why PostgreSQL caps `BLCKSZ` at 32 KB. This engine's 16/14/2 split gives the same practical range at 8 KB but is a different physical encoding.


PostgreSQL also attaches meaning to the flag states that this engine does not yet enforce: `LP_REDIRECT` and `LP_UNUSED` must have `lp_len == 0`, and `LP_DEAD` is set by index scans (the `kill_prior_tuple` optimisation) as well as by pruning.

3. **Free Space Allocation Formula:** $\text{Free Space Available} = \text{pd_upper} - \text{pd_lower} - 4$ Every tuple insertion requires both tuple payload memory ($\geq \text{len}$) and a new 4-byte line pointer in the header array.

PostgreSQL's `PageGetFreeSpace()` computes the same difference, but two extra rules apply that this engine does not implement:

- Every item is **MAXALIGNED** (8 bytes on 64-bit), so the true cost of a tuple is `MAXALIGN(len) + 4`, not `len + 4`.
 - `PageGetHeapFreeSpace()` additionally reports zero once the page already holds `MaxHeapTuplesPerPage` (291) line pointers, because no further line pointer can be addressed.
-

3. Tuple Insertion Algorithm



4. PostgreSQL Fidelity Check

Claim	Verdict
Slotted layout, header + line pointers growing up, tuples growing down	Exact
<code>pd_lsn</code> , <code>pd_checksum</code> , <code>pd_flags</code> , <code>pd_lower</code> , <code>pd_upper</code> , <code>pd_special</code> order and widths	Exact
<code>PD_ALL_VISIBLE = 0x0004</code>	Exact
Line pointer = 4 bytes; states <code>UNUSED/NORMAL/REDIRECT/DEAD</code> = 0/1/2/3	Exact
Page header is 18 bytes	Simplified — PostgreSQL is 24 (<code>pd_pagesize_version</code> , <code>pd_prune_xid</code> omitted)
Line-pointer bit split 16/14/2	Divergent — PostgreSQL is 15/2/15
<code>free = pd_upper - pd_lower - 4</code>	Simplified — PostgreSQL MAXALIGNs items and caps line pointers at 291
Page checksum	Not implemented — <code>pd_checksum</code> is stored but never computed; PostgreSQL uses an FNV-1a-based 16-bit page checksum

Implementation status (2026-08-27)

`Page` is now a **view** over 8192 bytes it does not own, so callers work directly inside a pinned buffer frame the way `BufferGetPage()` does. `PageBuffer` owns storage where a page genuinely needs its own (recovery scratch, fresh pages, tests).

`insert_tuple` no longer recycles `LP_DEAD` line pointers — only `LP_UNUSED` . A DEAD pointer may still be the target of an index entry VACUUM has not cleaned yet, and handing that slot to a new tuple makes the stale entry resolve to an unrelated row. `Page::insert_tuple_at` was added for WAL redo, which must restore a tuple to the exact slot the record names rather than letting the page choose.

4. Item 3: Tuples & Heap CTID Physical Addressing


Confidence: verified

Citations: [include/pg/tuple.h:1-90](#), [include/pg/heap.h:1-82](#), [src/heap.cpp:1-210](#), [tests/test_heap.cpp:1-120](#)

1. Physical CTID Tuple Addressing

Every row version in the database is uniquely addressed by its **CTID** — in PostgreSQL an `ItemPointerData`, surfaced to SQL as the system column `ctid`:

$\text{CTID} = (\text{page_id}, \text{slot_id})$



2. Invariants & Binary Layout

1. **TupleHeader Structure (16 Bytes — PostgreSQL uses 23)** (`[include/pg/tuple.h:42-47]`):

- `xmin` (4 Bytes): Transaction ID that inserted this row version.
- `xmax` (4 Bytes): Transaction ID that deleted or updated this row version (`0` if currently live).
- `t_ctid` (6 Bytes): Forward pointer `CTID {page, slot}` to the newest row version (or self-reference (page, slot) if newest). **Exact match** for PostgreSQL's `ItemPointerData` (4-byte block id + 2-byte offset).
- `infomask` (2 Bytes): Status flags — see the correction in invariant 2.

PostgreSQL's `HeapTupleHeaderData` is **23 bytes** before the null bitmap, MAXALIGNED to 24 in practice:

Field	Size	Present here?
<code>t_choice</code> union — <code>t_xmin</code> , <code>t_xmax</code> , then <code>t_cid</code> / <code>t_xvac</code>	12 B	partly (no <code>t_cid</code>)

Field	Size	Present here?
t_ctid	6 B	yes
t_infomask2	2 B	no
t_infomask	2 B	merged into one 16-bit word
t_hoff	1 B	no
t_bits[] null bitmap	variable	no

t_ctid carries correctness this engine does not reach: it is the *command id*, which is what stops a statement from seeing rows it inserted itself within the same transaction.

2. **Infomask correction.** This engine packs three flags into one infomask word. In PostgreSQL those flags live in two different words, and one of the three values is not PostgreSQL's:

Flag	cpp-postgres	PostgreSQL	PostgreSQL field
HEAP_HASEXTERNAL	0x2000	0x0004	t_infomask
HEAP_HOT_UPDATED	0x4000	0x4000 ✓	t_infomask2
HEAP_ONLY_TUPLE	0x8000	0x8000 ✓	t_infomask2

0x2000 in PostgreSQL's t_infomask is HEAP_UPDATED ("this tuple is the product of an UPDATE"), an unrelated flag. The values chosen here are internally consistent but are not PostgreSQL's bit assignments.

PostgreSQL additionally uses t_infomask for the **hint bits** HEAP_XMIN_COMMITTED (0x0100), HEAP_XMIN_INVALID (0x0200), HEAP_XMAX_COMMITTED (0x0400), HEAP_XMAX_INVALID (0x0800) — a per-tuple cache of the CLOG answer, so a repeated visibility test costs zero CLOG page reads. This engine has no hint bits and consults CLOG every time (Item 13).

3. **Physical Tuple Capacity Derivation:** For an 8KB page with 18-byte page header and 24-byte serialized heap tuples (16B header + 8B data item_id + price), the maximum row capacity per page is:

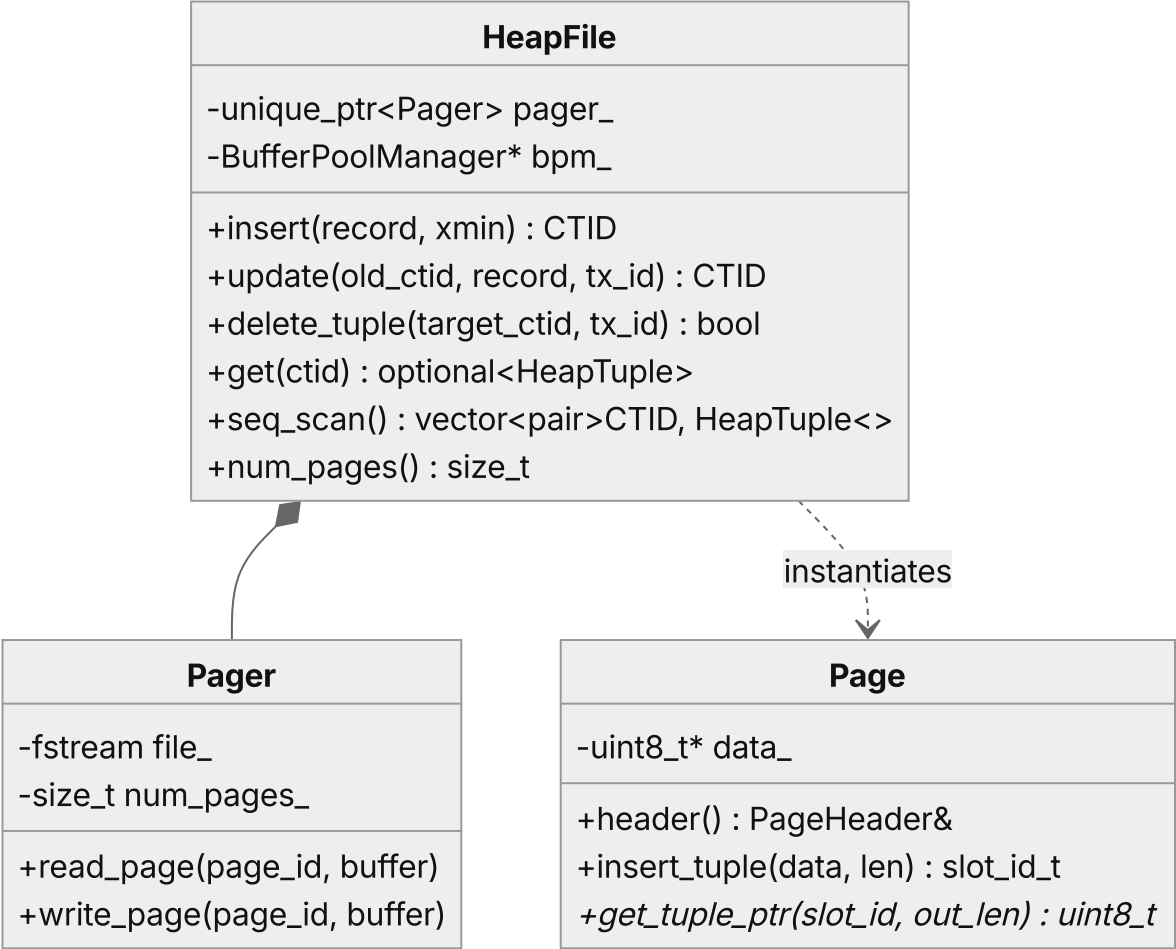
$$\text{Max Tuples Per Page} = \left\lfloor \frac{8192 - 18}{24 + 4} \right\rfloor = \left\lfloor \frac{8174}{28} \right\rfloor = 291 \text{ tuples}$$

By coincidence this equals PostgreSQL's MaxHeapTuplesPerPage, which reaches 291 by different arithmetic:

$$\left\lfloor \frac{8192 - 24}{\text{MAXALIGN}(23) + 4} \right\rfloor = \left\lfloor \frac{8168}{28} \right\rfloor = 291$$

In PostgreSQL 291 is a hard structural ceiling on line pointers per heap page, not merely the capacity for one particular row width.

3. Class Diagram: HeapFile & Pager Relationship



4. PostgreSQL Fidelity Check

Claim	Verdict
CTID = (block, offset), 6 bytes, offsets 1-based	Exact (<code>ItemPointerData</code>)
<code>t_ctid</code> self-references on the newest version, forward-points on update	Exact
<code>xmin</code> / <code>xmax</code> are 4-byte XIDs at the head of the tuple header	Exact
Tuple header is 16 bytes	Simplified — PostgreSQL is 23 B (<code>t_cid</code> , <code>t_infomask2</code> , <code>t_hoff</code> , <code>t_bits</code> omitted)
<code>HEAP_HASEXTERNAL</code> = <code>0x2000</code>	Wrong for PostgreSQL — real value is <code>0x0004</code> ; <code>0x2000</code> is <code>HEAP_UPDATED</code>
<code>HEAP_HOT_UPDATED</code> / <code>HEAP_ONLY_TUPLE</code> live in <code>infomask</code>	Wrong field — PostgreSQL stores both in <code>t_infomask2</code>
291 tuples per page	Coincidentally exact — matches <code>MaxHeapTuplesPerPage</code>

Claim	Verdict
<code>xmax == 0</code> means live	Simplified — PostgreSQL also writes <code>xmax</code> for row <i>locks</i> (<code>HEAP_XMAX_LOCK_ONLY</code>) and for <code>MultiXactId</code> s, neither of which deletes the row

5. Item 4: MVCC & Transaction Snapshot Visibility

Confidence: verified


Citations: [include/pg/tx.h:1-75](#), [include/pg/mvcc.h:1-45](#), [src/tx.cpp:1-95](#), [src/mvcc.cpp:1-75](#), [tests/test_mvcc.cpp:1-135](#)

1. Multi-Version Concurrency Control (MVCC) Overview

PostgreSQL implements multi-version concurrency control using row-version header stamps (`xmin` and `xmax`). Readers never block writers, and writers never block readers.

Two clarifications the rest of this chapter depends on:

- **Snapshot Isolation is not the default.** PostgreSQL's default isolation level is **Read Committed**, which takes a *fresh snapshot at the start of every statement*. Snapshot Isolation in the textbook sense (one snapshot for the whole transaction) is PostgreSQL's `REPEATABLE READ`. `SERIALIZABLE` layers Serializable Snapshot Isolation (SSI) predicate-lock tracking on top of that. `cpp-postgres` implements exactly one model — one snapshot per transaction — which corresponds to `REPEATABLE READ`.
- **A snapshot's `xmin` does not mean "everything below is committed."** It means everything below is *finished* — committed **or** aborted. Whether a sub-`xmin` transaction committed still has to be answered from CLOG (Item 13). PostgreSQL's own comment on `SnapshotData` is that XIDs `< xmin` are "visible to the snapshot" only in the sense that no further in-progress test is needed.



2. Invariants & Visibility State Machine

A tuple version is visible to a `Snapshot` if and only if **both** of the following conditions hold (`[src/mvcc.cpp:25]`):

0. **Preliminary:** this engine evaluates every `xmin` / `xmax` by calling into CLOG. PostgreSQL first checks the tuple's **hint bits** (`HEAP_XMIN_COMMITTED`, `HEAP_XMAX_INVALID`, ...) and only falls back to CLOG on a miss, then writes the answer back into the tuple header. That write is why a pure `SELECT` can dirty pages in PostgreSQL — a behaviour this engine does not exhibit.

1. Creation Visibility (`xmin`):

- `xmin == snapshot.current_tx_id` (created by current tx), OR
- (`xmin < snapshot.xmax` AND `xmin` is NOT in `snapshot.active_txs` AND `status(xmin) == COMMITTED`).

2. Deletion Invisibility / Retention (`xmax`): A tuple remains visible regarding its deletion state if:


- `xmax == 0` (never deleted or updated), OR
- `status(xmax) == ABORTED` (deleter transaction rolled back), OR
- `status(xmax) == IN_PROGRESS` (deleter transaction is still running concurrently), OR
- `xmax >= snapshot.xmax` (deleted in the future after this snapshot was taken), OR
- `xmax` IS in `snapshot.active_txs` (deleter was in-flight when snapshot started). (Note: If `xmax == snapshot.current_tx_id`, the deletion was executed by the current transaction itself, making the tuple invisible).

Where PostgreSQL is harder than this. The rule above assumes `xmax != 0` means "deleted or updated". In PostgreSQL `xmax` is overloaded:

- `HEAP_XMAX_LOCK_ONLY (0x0080)` — `xmax` records a row **lock** (`SELECT ... FOR UPDATE/SHARE`), not a deletion. The row is still live.
- `HEAP_XMAX_IS_MULTI (0x1000)` — `xmax` is a `MultiXactId`, a handle for *several* transactions locking the same row, resolved through the `pg_multixact` SLRUs.
- `t_cid` and `HEAP_COMBOCID (0x0020)` — within one transaction, a row inserted by command 1 must be invisible to command 1 itself but visible to command 2. This engine has no command ids, so intra-transaction statement ordering is not modelled.

PostgreSQL's real implementation lives in `HeapTupleSatisfiesMVCC()` in `heapam_visibility.c`, and it is roughly 150 lines of exactly these special cases.

3. Sequence Diagram: Concurrent MVCC Reads and Writes



4. PostgreSQL Fidelity Check

Claim	Verdict
MVCC via <code>xmin</code> / <code>xmax</code> stamps; readers do not block writers	Exact
Snapshot = (xmin, xmax, active_txs) triple	Exact — PostgreSQL's <code>SnapshotData</code> is <code>xmin</code> , <code>xmax</code> , <code>xip[]</code>
Visibility = "created by a committed, non-concurrent tx AND not deleted by one"	Exact in shape
"PostgreSQL implements Snapshot Isolation"	Corrected — that is <code>REPEATABLE READ</code> ; the default is Read Committed, one snapshot <i>per statement</i>
"All Tx < snapshot.xmin are COMMITTED"	Corrected — they are <i>finished</i> ; committed vs aborted still comes from CLOG
CLOG consulted on every visibility test	Simplified — PostgreSQL short-circuits through hint bits first
<code>xmax != 0</code> implies deleted/updated	Simplified — PostgreSQL also uses <code>xmax</code> for row locks and MultiXactIds
No command-id (<code>t_cid</code>) handling	Missing — affects visibility <i>within</i> a single transaction
No XID freezing / wraparound handling	Missing — PostgreSQL must freeze <code>xmin</code> before the 32-bit XID space wraps, or refuse writes

Implementation status (2026-08-27)

Transaction state now lives in a `Session`, one per connection, so two clients no longer share a transaction and two snapshots can coexist — which is the situation these visibility rules exist for. `TransactionManager` publishes each running transaction's snapshot xmin (PostgreSQL's `PGPROC.xmin`) so the VACUUM horizon is computed from snapshots rather than transaction ids.

A `LockManager` serialises writers on the same row. MVCC keeps readers out of the way, but two transactions updating one row must not both succeed; with a single-threaded executor there is nobody to wait for, so the conflict is reported rather than silently losing the first write.

Still missing: command ids (`t_cid`), hint bits, MultiXacts, XID freezing.

6. Item 5: VACUUM Engine & In-Place Page Defragmentation


Confidence: verified

Citations: [include/pg/vacuum.h:1-35](#), [src/vacuum.cpp:1-120](#), [tests/test_vacuum.cpp:1-115](#)

1. Dead Tuple Reclamation Mechanics

When rows are updated or deleted under MVCC, old tuple versions become **dead** once their `xmax` falls below the oldest snapshot that could still see them — PostgreSQL's `OldestXmin`, computed across every backend and replication slot in the cluster:

`(\text{Dead Condition: } (\text{xmax} > 0) \wedge (\text{status}(\text{xmax}) = \text{COMMITTED}) \wedge (\text{xmax} < \text{oldest\_active\_xmin}))`

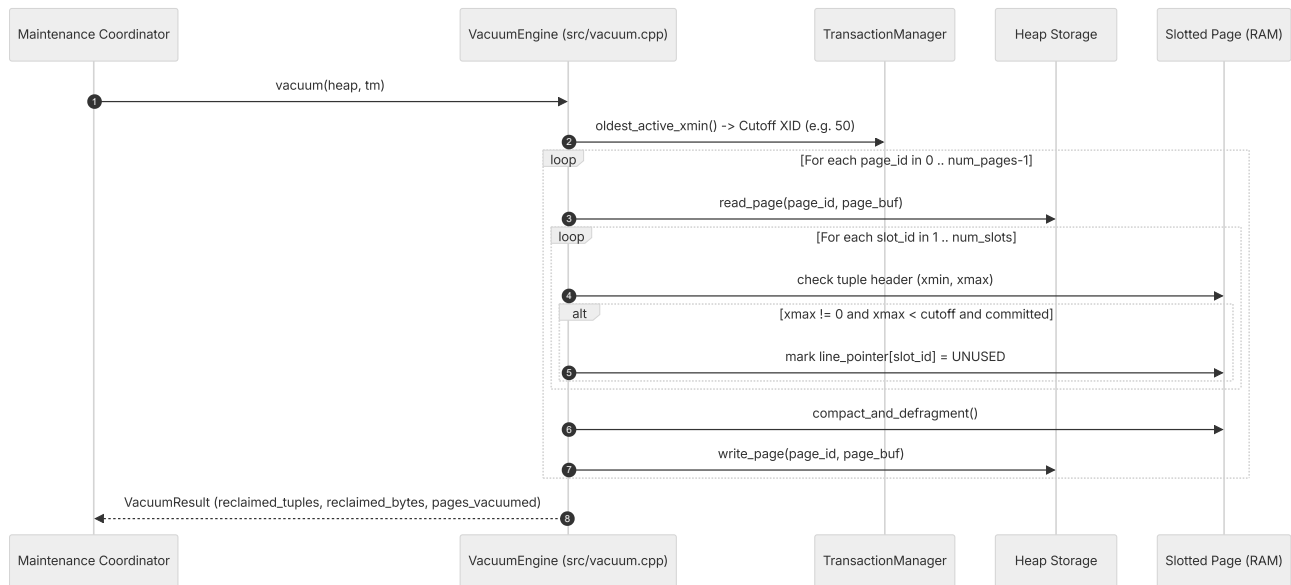


2. Invariants & CTID Stability

- 1. **CTID Slot Stability:** VACUUM never shifts or rennumbers existing `slot_id` line pointer indexes. Slot 3 remains `(page, 3)`, guaranteeing that foreign references and secondary index entries are not corrupted (`[src/page.cpp:125]`).

2. **Compact Contiguous Allocation:** Surviving tuple payloads are slid upward towards the end of the 8KB page (`PAGE_SIZE = 8192`), coalescing all freed memory into a single contiguous block between `pd_lower` and `pd_upper` .

3. Sequence Diagram: Vacuum Page Lifecycle




4. PostgreSQL Fidelity Check

4.1 VACUUM is three phases in PostgreSQL, not one

The single-pass loop diagrammed above — *see a dead tuple, mark its line pointer `UNUSED`, defragment* — is only correct for a **table with no indexes**. As soon as an index exists, marking a slot `UNUSED` immediately would leave index entries pointing at a slot that a later insert can reuse, and the index would silently return the wrong row.

PostgreSQL therefore splits the work:



`LP_DEAD` -> `LP_UNUSED` only becomes legal *after* every index has been cleaned. When the dead-TID store fills up, PostgreSQL runs phases 2 and 3 and then resumes phase 1 — a large VACUUM can cycle several times.

4.2 What else PostgreSQL's VACUUM does

- **Freezing.** The primary reason VACUUM is mandatory, not optional. XIDs are 32-bit and wrap; VACUUM rewrites old `xmin` values as frozen (via `HEAP_XMIN_FROZEN` hint bits) so that ancient rows stay visible after wraparound. Left unrun long enough, PostgreSQL refuses new writes. **Not modelled here at all.**
- **Visibility map.** Two bits per heap page (`all-visible` , `all-frozen`) maintained by VACUUM. This is what makes index-only scans and page-skipping possible. **Not modelled.**
- **Free space map.** A separate `_fsm` fork recording per-page free space so inserts can find a home without scanning. **Not modelled** — this engine appends or scans linearly.
- **Truncation.** VACUUM can return trailing all-empty pages to the OS; otherwise space is reused inside the relation, never handed back. **Not modelled.**
- **HOT pruning happens outside VACUUM.** `heap_page_prune_opt()` runs during ordinary page access, guarded by the `pd_prune_xid` hint (Item 2). Most dead-tuple cleanup in a healthy PostgreSQL is done by SELECTs, not by VACUUM.
- **VACUUM vs VACUUM FULL** . Plain VACUUM defragments in place, exactly as modelled here. `VACUUM FULL` rewrites the entire relation into a new file and takes an `ACCESS EXCLUSIVE` lock — a different algorithm.
- **Autovacuum.** A launcher plus worker processes driven by dead-tuple thresholds. This engine vacuums only on explicit command.

4.3 Verdict table

Claim	Verdict
Dead = <code>xmax</code> committed and below the global oldest xmin	Exact in principle (<code>OldestXmin</code>)

Claim	Verdict
VACUUM never rennumbers slot ids, preserving CTID stability for indexes	Exact — this is the core invariant PostgreSQL relies on
Surviving tuples slide toward the page end, coalescing free space	Exact — <code>PageRepairFragmentation()</code>
One pass, dead tuple straight to <code>UNUSED</code>	Divergent — safe only without indexes; PostgreSQL needs <code>LP_DEAD</code> -> index vacuum -> <code>LP_UNUSED</code>
Freezing / wraparound, visibility map, FSM, truncation, autovacuum	Missing

Implementation status (2026-08-27)

VACUUM is now the three-phase algorithm described above, not a single pass:

1. scan the heap, flag dead tuples `LP_DEAD`, collect their TIDs
2. remove every index entry pointing at a collected TID
3. demote `LP_DEAD` to `LP_UNUSED`, then compact

A dead HOT chain root becomes `LP_REDIRECT` rather than being freed, so the index entry that points at it keeps working while the tuple bytes go away — and index scans follow the redirect via

```
HeapFile::hot_search .
```

It also runs through the buffer pool. Reading the relation directly while every other subsystem went through the pool meant VACUUM acted on stale pages and its work was later overwritten by whichever cached frame flushed last.

Still missing: freezing and wraparound handling, the visibility map, the free space map fork, relation truncation, autovacuum.

7. Item 6: Secondary B-Tree Index Subsystem

Confidence: verified


Citations: [include/pg/btree.h:1-69](#), [src/btree.cpp:1-85](#), [tests/test_index.cpp:1-125](#)

1. Index Decoupling & Multi-Version Addressing

In PostgreSQL, **every** index is a secondary index — there is no clustered/primary storage — and an index entry maps a key to a heap CTID. Critically, an index tuple carries **no MVCC header**: no `xmin`, no `xmax`. This is the property that forces the heap dereference in §2.

Two corrections to the stronger form of that claim:

- Index tuples **do** have a header. PostgreSQL's `IndexTupleData` is 8 bytes: `t_tid` (6 B, the heap CTID) plus `t_info` (2 B, holding the tuple size, a `INDEX_VAR_MASK` bit, and `INDEX_NULL_MASK`). When `INDEX_NULL_MASK` is set, a **null bitmap follows** — so index tuples do store null bitmaps, they just do not store visibility information.
- Index tuples **can** store extra columns. `CREATE INDEX ... INCLUDE (cols)` (covering indexes, v11+) stores non-key payload columns in the leaf tuples specifically so that index-only scans can answer without touching the heap.




2. Invariants & Lookup Pipeline

- Multi-Version Key Duplication:** Because updates create new row versions at new physical CTIDs without deleting the old version, a single index key can map to multiple candidate CTIDs simultaneously (`[src/btree.cpp:22]`).

2. **Mandatory Heap MVCC Dereference:** The index lookup returns candidate CTIDs. The database must dereference each CTID against the physical Heap and evaluate snapshot visibility rules before returning tuples to the query caller ([src/engine.cpp:140]).

3. Sequence Diagram: Index Point Scan



4. PostgreSQL Fidelity Check

Claim	Verdict
Index maps key -> CTID; the heap is the only source of row data	Exact
Index tuples carry no <code>xmin</code> / <code>xmax</code> , so every hit needs a heap visibility check	Exact — the defining property of PostgreSQL's index design
One key can map to several CTIDs because updates create new versions	Exact
"Index does not store null bitmaps"	Corrected — <code>IndexTupleData.t_info</code> has <code>INDEX_NULL_MASK</code> and a null bitmap follows when set
"Index does not store row columns"	Corrected for modern PostgreSQL — <code>INCLUDE</code> columns are stored in leaf tuples

Mechanisms PostgreSQL adds that this engine does not model

- **Index-only scans.** If the visibility map says every tuple on the target heap page is all-visible, PostgreSQL skips the heap fetch entirely and answers from the index tuple. Without a visibility map (Item 5) this engine can never do that — the heap dereference in §2 is genuinely mandatory *here*, but only usually mandatory in PostgreSQL.
- `kill_prior_tuple` / `LP_DEAD` **index hints.** When an index scan dereferences a CTID and finds the tuple dead to everyone, it flags that index entry `LP_DEAD` so later scans skip it without a heap read, and so the page can be cleaned without a full index vacuum.
- **B-tree deduplication (v13+).** Repeated keys are folded into a single *posting list* tuple holding a sorted TID array, which dramatically shrinks indexes on low-cardinality columns. In this engine every duplicate key costs a full entry.
- **Bitmap index scans.** For a wide match set the planner collects TIDs into a bitmap, sorts by physical page, and scans the heap in page order to convert random I/O into sequential I/O. This engine always does one random heap fetch per candidate.
- **Other access methods.** GiST, GIN, SP-GiST, BRIN and hash all exist behind the same `Key -> CTID` contract.

8. Item 7: Heap-Only Tuples (HOT) & Update Chains

Confidence: verified

Citations: [include/pg/heap.h:40-50](#), [src/heap.cpp:115-180](#), [tests/test_hot.cpp:1-130](#)


1. The Write-Amplification Problem & HOT Solution

In standard MVCC, updating an unindexed column (e.g. `UPDATE items SET price = 20 WHERE item_id = 100;`) forces a new entry into every secondary index on the table, leading to severe **index bloat**.

Heap-Only Tuples (HOT) eliminates all secondary index writes when:

1. No indexed column of **any** index on the table is modified by the update.
2. The new tuple version fits onto the **same 8KB page** as the old version.


Both conditions are PostgreSQL's, exactly. Condition 2 is why `fillfactor` matters: lowering it below 100 reserves free space on each page at load time specifically to keep future updates HOT-eligible. Condition 1 is evaluated against the *union* of all indexed columns, which is why adding one index on a hot-updated column can collapse HOT for the whole table.



2. Invariants & Infomask Flags

1. **HEAP_HOT_UPDATED** (`0x4000`): Stamped on the old tuple version. Indicates that the next version in the update chain resides on this exact same page (`[include/pg/tuple.h:51]`). *The value matches PostgreSQL, but PostgreSQL stores it in `t_infomask2`, not `t_infomask` — this engine has only one infomask word (Item 3).*
2. **HEAP_ONLY_TUPLE** (`0x8000`): Stamped on the newly created tuple version. Tells the engine that no index pointer references this tuple directly; it can only be reached by following the HOT chain from the root line pointer (`[include/pg/tuple.h:52]`). *Same note: `t_infomask2` in PostgreSQL.*
3. **ItemFlags::REDIRECT** : When pruning removes the dead root tuple's data, it converts Slot 1's line pointer into a redirect (`lp_flags = REDIRECT, lp_offset = slot_2`), preserving index stability with zero tuple bytes. **This is exactly PostgreSQL's `LP_REDIRECT`**, and it is the reason the root line pointer of a HOT chain can never be reused: the index still points at it.
4. **Pruning is not VACUUM**. In PostgreSQL, HOT chain collapsing is done by `heap_page_prune_opt()` during ordinary page access — including plain `SELECT` s — guarded by the `pd_prune_xid` page-header hint (Item 2). `VACUUM` also prunes, but it is not the primary mechanism. Because this engine has no `pd_prune_xid` and no opportunistic prune path, HOT chains here are only collapsed by an explicit `VACUUM`.

3. Sequence Diagram: HOT Chain Traversal



4. PostgreSQL Fidelity Check

Claim	Verdict
HOT applies when no indexed column changes and the new version fits on the same page	Exact
Old version gets <code>HEAP_HOT_UPDATED</code> , new version gets <code>HEAP_ONLY_TUPLE</code>	Exact (values <code>0x4000</code> / <code>0x8000</code> are PostgreSQL's)
<code>t_ctid</code> links the chain forward within the page	Exact
Index keeps pointing only at the root line pointer	Exact
Pruning converts the root line pointer to <code>LP_REDIRECT</code>	Exact
Those flags live in <code>infomask</code>	Wrong field — PostgreSQL puts both in <code>t_infomask2</code>
Pruning happens in VACUUM	Incomplete — PostgreSQL prunes opportunistically on page access via <code>heap_page_prune_opt()</code>
No <code>fillfactor</code> control	Missing — the main knob DBAs use to keep updates HOT
HOT chain cannot span pages	Exact — an update that must move to another page falls back to a normal update plus index insert

9. Item 8: Shared Buffers & Clock-Sweep Replacement

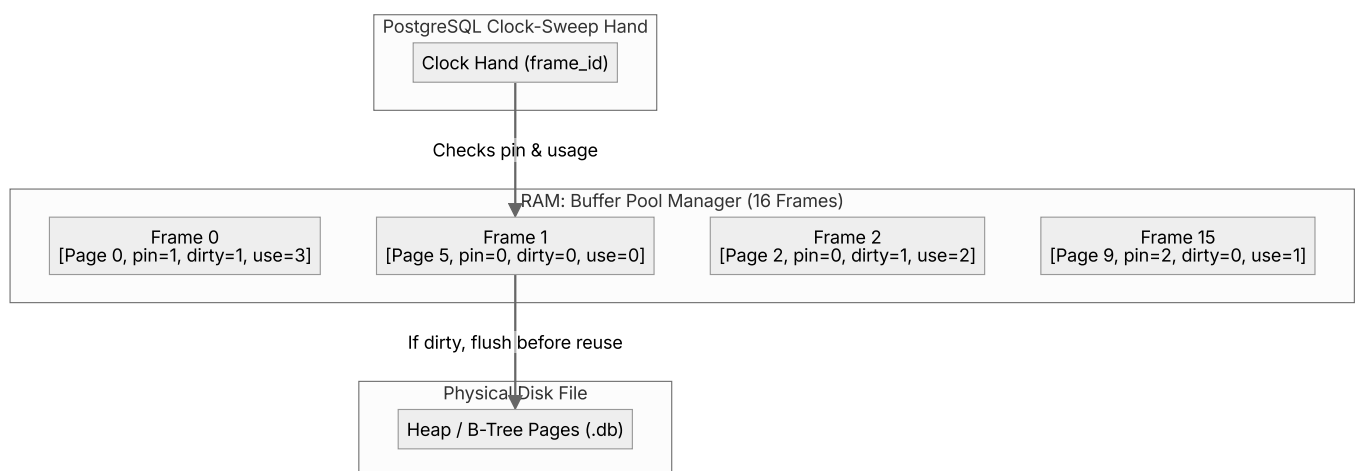
Confidence: `verified`

Citations: `include/pg/buffer_pool.h:1-75`, `src/buffer_pool.cpp:1-125`, `tests/test_buffer_pool.cpp:1-120`

1. Buffer Pool Manager (Shared Buffers) Architecture

`BufferPoolManager` maintains a configurable pool of in-memory frames (defaulting to **16 frames**, each exactly 8,192 bytes) acting as an LRU-approximation cache between the execution engine and physical disk files.

PostgreSQL's equivalent is `shared_buffers`, defaulting to **128 MB = 16,384 buffers**. Sixteen frames here is a teaching size chosen so that eviction is observable in the GUI after a handful of inserts.



2. Invariants & Clock-Sweep State Machine

1. **Pin Count Protection** (`pin_count > 0`): Any frame currently pinned by an active query worker is immediately skipped by the clock-sweep hand; it can **never** be evicted while in use (`[src/buffer_pool.cpp:85]`).

2. **Second-Chance Algorithm** (`usage_count`):

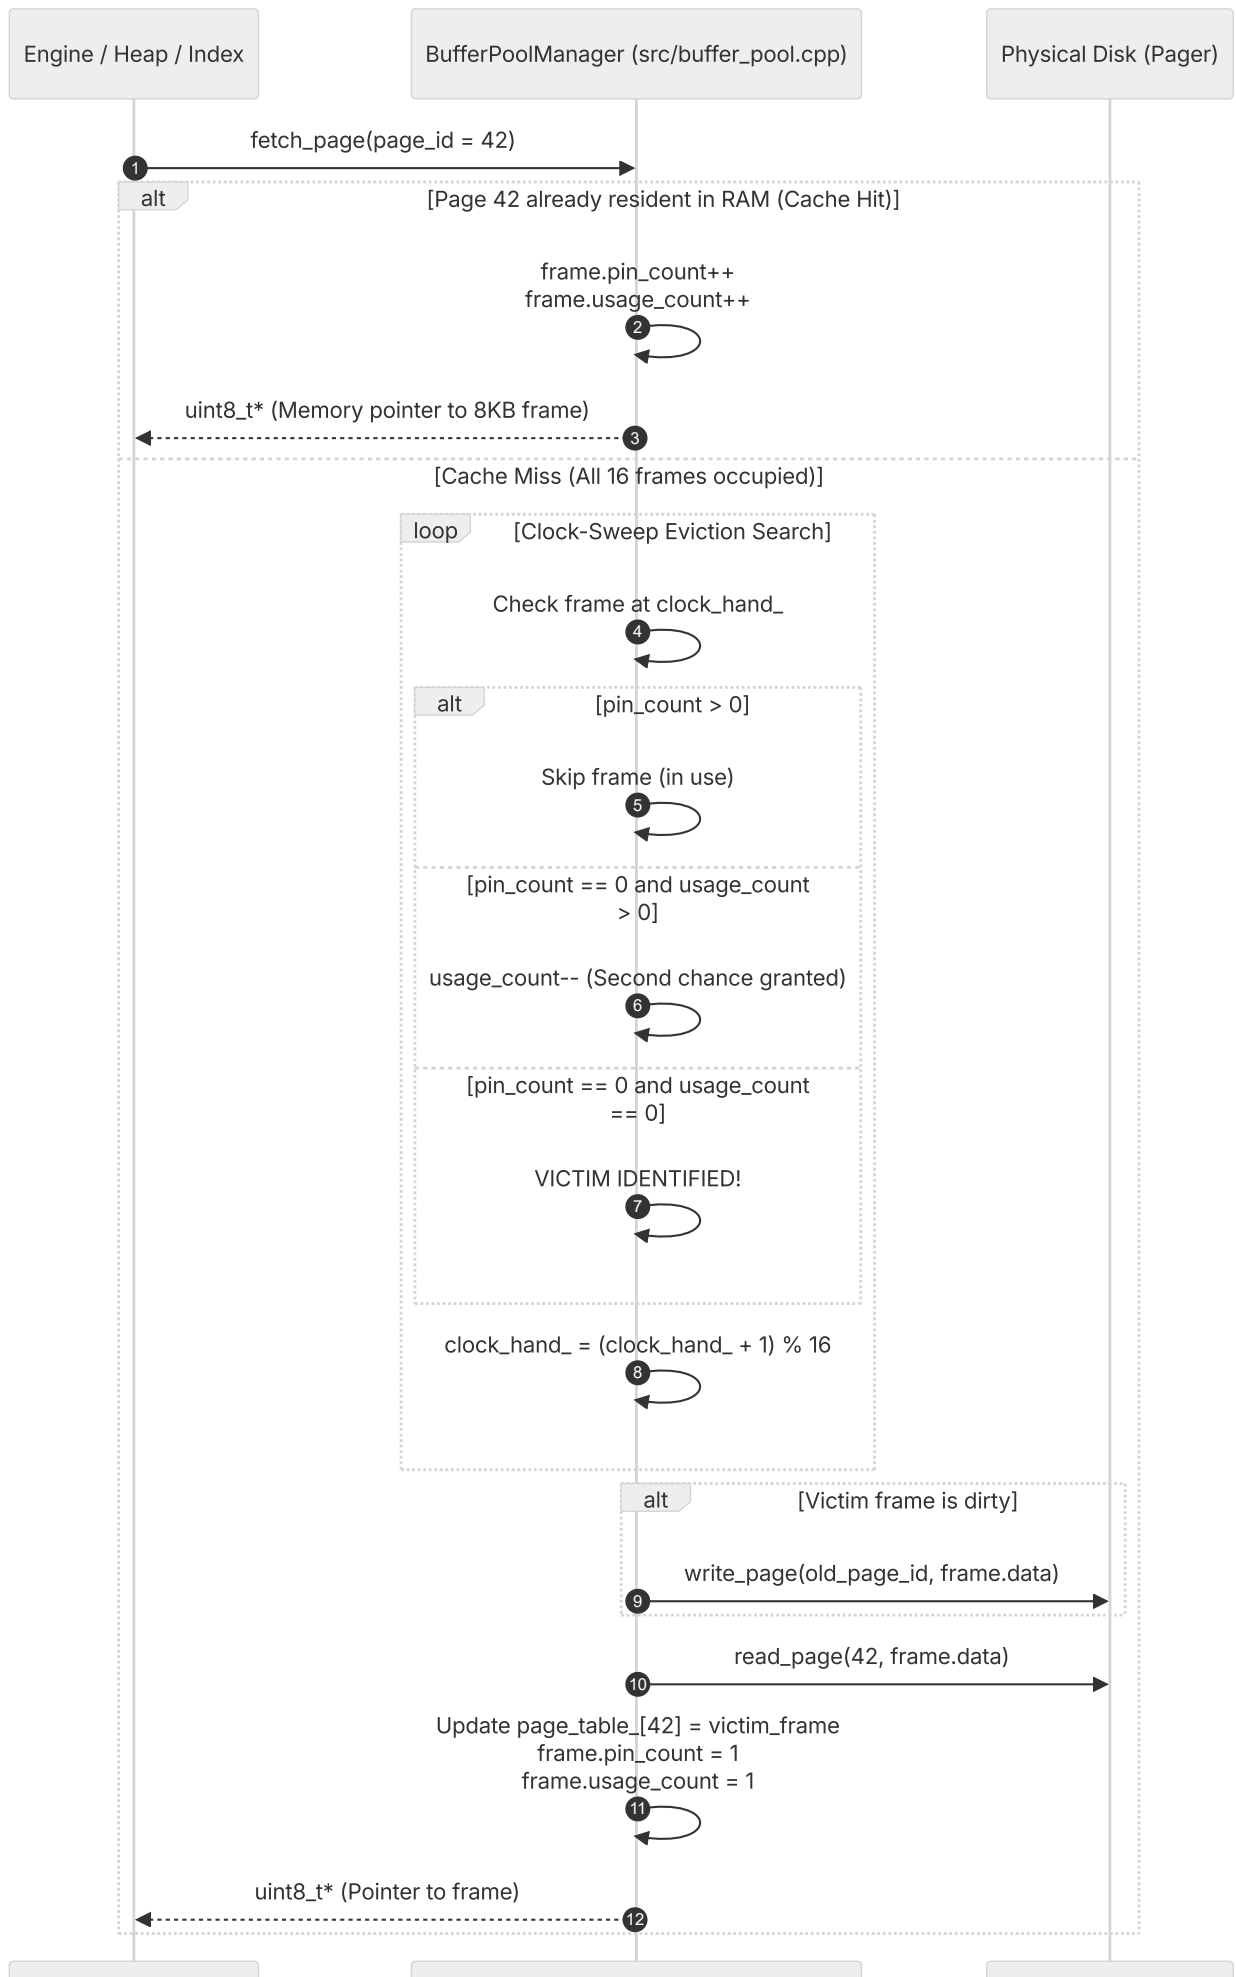
- If `pin_count == 0` and `usage_count > 0`, the clock hand decrements `usage_count--` and advances to the next frame.
- If `pin_count == 0` and `usage_count == 0`, the frame is selected as the eviction victim.

PostgreSQL's `StrategyGetBuffer()` is the same algorithm with one bound this engine lacks: `usage_count` saturates at `BUF_USAGECOUNT_MAX = 5`. Without a cap, a page touched ten thousand times would need ten thousand hand passes to become evictable, and the sweep degenerates. PostgreSQL also consults a **freelist** of never-used or explicitly freed buffers *before* starting the sweep.

3. **Dirty Writeback Before Eviction**: If the victim frame has `is_dirty == true`, it is written to disk via `Pager::write_page()` before being repurposed for the new page.

4. **WAL Ordering Applies to Eviction Too:** PostgreSQL's `FlushBuffer()` calls `XLogFlush(page.pd_lsn)` *before* the write, so evicting a dirty buffer can force a WAL flush. This is the same invariant as Item 9, enforced at the eviction site.

3. Sequence Diagram: Page Fetch & Eviction Lifecycle



4. PostgreSQL Fidelity Check

Claim	Verdict
Fixed array of 8KB frames, hash table from page id to frame	Exact (<code>BufferDesc</code> array + <code>SharedBufHash</code>)
Clock-sweep with pin count and usage count	Exact (<code>StrategyGetBuffer</code>)
Pinned buffers are never evicted	Exact
Dirty victim is written before reuse	Exact
16 frames	Teaching size — PostgreSQL defaults to 16,384 (128 MB)
<code>usage_count</code> unbounded	Divergent — PostgreSQL caps at <code>BUF_USAGECOUNT_MAX = 5</code>
Sweep starts immediately on miss	Simplified — PostgreSQL checks the freelist first
No WAL flush on eviction	Divergent — PostgreSQL's <code>FlushBuffer()</code> calls <code>XLogFlush()</code> first

Not modelled

- **Buffer access strategies (ring buffers).** Sequential scans of large tables, `COPY`, `CREATE TABLE AS` and `VACUUM` each run inside a small ring (`BufferAccessStrategy`) so a single big scan cannot flush the whole cache. Without this, one large scan in PostgreSQL would evict everything.
- **The background writer.** A dedicated process that trickles dirty buffers out ahead of the clock hand so foreground backends rarely have to write during eviction.
- **Local buffers.** Temporary tables use a per-backend local buffer pool, not shared buffers.
- **Buffer content locks and pin/refcount separation.** PostgreSQL splits a *pin* (the page will not be evicted) from a *content lock* (shared/exclusive access to the bytes), and tracks both in shared memory with atomic state words. This engine has a single-mutex model.
- **Double buffering.** PostgreSQL deliberately keeps `shared_buffers` modest and relies on the OS page cache underneath; sizing advice follows from that layering.

Implementation status (2026-08-27)

Three changes bring the pool closer to the real thing:

- `usage_count` is capped at `BUF_USAGECOUNT_MAX = 5`. Uncapped, a hot page needed as many hand passes as it had hits, and the sweep gave up while frames sat unpinned and evictable — it threw "all frames are pinned" when none were. The cap is what makes the sweep terminate.
- **The WAL rule is enforced at eviction.** `write_frame` calls `WALManager::flush_up_to(page.lsn())` before any page reaches disk.
- **Pages are pinned across read-modify-write.** `PinnedPage` is an RAIL pin, and `fetch_page` returns a view into the frame rather than a copy, so a caller mutates shared memory instead of a private copy it later writes back over whatever happened in between.

Still missing: strategy rings, a background writer, local buffers for temp relations, and separate pin/content-lock states.


10. Item 9: Write-Ahead Logging (WAL) & REDO Durability

Confidence: verified

Citations: [include/pg/wal.h:1-95](#), [src/wal.cpp:1-180](#), [tests/test_wal.cpp:1-135](#)

1. Write-Ahead Logging (WAL) Architecture

To guarantee **ACID Durability** and high write throughput, PostgreSQL converts random 8KB page writes into sequential, append-only log records in the Write-Ahead Log (`*.log`).




2. Invariants & Record Binary Layout

1. **The WAL Invariant:** `page.pd_lsn` $\leq$ `wal.flushed_lsn` A dirty page in RAM is never written to disk before all WAL records up to the page's `pd_lsn` have been flushed to disk (`[src/wal.cpp:110]`).
2. **Log Record Structure (35-Byte Header + Dynamic Payload — this engine's own format)** (`[include/pg/wal.h:35-47]`):
 - `lsn` (8 Bytes): Monotonically increasing Log Sequence Number. *In PostgreSQL an LSN is not a counter — it is a byte position in the logical WAL stream, which is why `pg_current_wal_lsn()` prints as `hex/hex`.*

- `prev_lsn` (8 Bytes): Backward pointer for UNDO / transaction rollback chains. **This has no PostgreSQL analogue.** PostgreSQL's `xl_prev` links to the *immediately preceding record in the stream* and exists to detect stale/recycled WAL pages during recovery, not to drive rollback — PostgreSQL never rolls back from WAL (Item 16).
- `tx_id` (4 Bytes): Transaction ID.
- `type` (1 Byte): `1` =INSERT, `2` =UPDATE, `3` =COMMIT, `4` =ABORT, `5` =CHECKPOINT, `6` =FPI, `7` =CLR.
- `page_id` (4 Bytes), `slot_id` (2 Bytes): Target heap page and line pointer.
- `payload_len` (4 Bytes): Byte length of following payload (`uint32_t`).
- `crc` (4 Bytes): CRC-32 (IEEE 802.3 polynomial) checksum verifying record integrity. *PostgreSQL has used **CRC-32C (Castagnoli)** since 9.5, chosen for the `SSE4.2` `crc32` instruction.*

PostgreSQL's `XLogRecord` header is **24 bytes**: `xl_tot_len` (4), `xl_xid` (4), `xl_prev` (8), `xl_info` (1), `xl_rmid` (1), 2 bytes padding, `xl_crc` (4). Everything after that is resource-manager-specific and self-describing: `xl_rmid` names the resource manager (heap, btree, xact, ...), and a chain of `XLogRecordBlockHeader` entries describes each page the record touches — including whether a full-page image is attached (Item 15). There is no fixed `page_id` / `slot_id` in PostgreSQL's header; a single record can modify several blocks in several relations.

3. Sequence Diagram: REDO Crash Recovery Pass



4. PostgreSQL Fidelity Check

Claim	Verdict
WAL turns random page writes into sequential log writes	Exact
<code>page.pd_lsn <= flushed_lsn</code> before a dirty page reaches disk	Exact — the WAL rule, enforced in <code>FlushBuffer()</code> / <code>XLogFlush()</code>

Claim	Verdict
Commit is acknowledged only after the WAL record is fsynced	Exact by default — <code>synchronous_commit = on</code> ; PostgreSQL lets you trade this away (<code>off</code> , <code>local</code> , <code>remote_write</code> , <code>remote_apply</code>)
CRC per record, torn tail detected by CRC mismatch and scan stops	Exact in shape
Redo replays inserts/updates for committed transactions	Exact in shape
CRC-32 IEEE	Divergent — PostgreSQL uses CRC-32C
35-byte fixed header with <code>page_id</code> + <code>slot_id</code>	Divergent — PostgreSQL: 24-byte <code>XLogRecord</code> + rmgr block headers, multi-block capable
<code>prev_lsn</code> used for undo chains	No analogue — PostgreSQL has no undo; <code>xl_prev</code> is a stream-integrity link
Recovery scans from byte 0	Divergent — PostgreSQL starts at the checkpoint's REDO pointer from <code>pg_control</code> (Item 15)
LSN as an abstract counter	Divergent — a PostgreSQL LSN is a byte offset in the WAL stream

Not modelled

- **WAL segmentation.** PostgreSQL writes 16 MB segment files under `pg_wal/` , recycles them, and archives them (`archive_command`). One flat `.log` file here.
- **Resource managers.** Heap, btree, xact, clog, gist, ... each register redo callbacks; `pg_waldump` decodes them.
- **WAL levels.** `minimal` / `replica` / `logical` change what gets logged; `wal_level = logical` is what makes logical decoding and CDC possible.
- **The WAL writer process**, group commit, and `commit_delay` .
- **Streaming replication.** The same WAL stream is shipped to standbys; `wal_sender` / `wal_receiver` , replication slots, and hot standby all ride on this file format.

Implementation status (2026-08-27)

The write-ahead rule is now actually enforced, in three places that all had to change together:

- **`log_commit` is durable.** `WALManager::flush` calls `File::sync()` , a real `fdatasync` / `_commit` . Previously it called `std::fstream::flush()` , which leaves the bytes in the OS page cache — a power cut lost committed work.
- **The record is written before the page changes.** `HeapFile` emits it while holding the page pinned, then applies the change, then stamps `pd_lsn` . The engine no longer logs after the fact, which was write-behind logging.
- **`pd_lsn` is set on every modification**, so the buffer pool can refuse to write a page out until the log covering it is durable.

Still divergent: CRC-32 rather than CRC-32C, a 35-byte fixed header rather than `XLogRecord` plus rmgr block references, one log file rather than 16MB segments.

11. Item 10: TOAST Oversized-Attribute Storage

Confidence: verified

Citations: [include/pg/toast.h:1-65](#), [src/toast.cpp:1-120](#), [tests/test_toast.cpp:1-115](#)


1. The TOAST Architecture

TOAST (The Oversized-Attribute Storage Technique) prevents large text, JSON, and binary attributes from degrading the row density of 8KB slotted heap pages.

When an attribute payload exceeds **2,048 bytes (2KB)**, the storage engine slices the payload into discrete 2,048-byte chunks stored in a dedicated auxiliary TOAST relation, replacing the inline data with a compact **18-byte `ToastPointer`**.

PostgreSQL's real numbers and real trigger. Both constants here are round approximations, and the trigger condition is different in kind:

- The threshold is `TOAST_TUPLE_THRESHOLD` = **2032 bytes**, and it is tested against the **size of the whole tuple**, not of one attribute. A row of six 500-byte columns gets toasted; a single 1 KB column in an otherwise narrow row does not.
- The chunk size is `TOAST_MAX_CHUNK_SIZE` = **1996 bytes**, sized so that exactly four chunk tuples fit in an 8 KB TOAST page after the heap tuple header and `chunk_id / chunk_seq` columns.
- Compression comes first.** PostgreSQL's `heap_toast_insert_or_update()` loops: compress the largest attribute (pglz, or lz4 when `default_toast_compression = lz4`), re-measure the tuple; if still over threshold, push the largest attribute out of line; repeat. Out-of-line storage is the *fallback*, not the first move. This engine chunks uncompressed bytes directly.
- Behaviour is per-column configurable via `ALTER TABLE ... SET STORAGE : PLAIN` (never toast), `EXTENDED` (compress then externalize — the default), `EXTERNAL` (externalize without compressing, so substring reads stay cheap), `MAIN` (compress, externalize only as a last resort).



2. Invariants & Binary Layout

1. `ToastPointer` Structure (18 Bytes) (`[include/pg/toast.h:28-36]`):

- `toast_id` (8 Bytes): Unique 64-bit identifier linking to the auxiliary relation (`uint64_t`).
- `raw_size` (4 Bytes): Original uncompressed byte size of the attribute (`uint32_t`).
- `chunk_count` (4 Bytes): Total number of auxiliary chunks (`\lceil \text{raw\_size} / 2048 \rceil`, `uint32_t`).

- `flags` (2 Bytes): Compression and storage strategy flags (`uint16_t`).

PostgreSQL's on-disk TOAST pointer is **also 18 bytes**, but the fields differ: a 1-byte `varattrib_1b_e` header, a 1-byte `va_tag`, then a 16-byte `varatt_external` of `va_rawsize` (int32), `va_extsize` (int32), `va_valueid` (Oid) and `va_toastrelid` (Oid). Two consequences of that layout:

- PostgreSQL stores the *relation* OID in the pointer, so a detoast can find its TOAST table from the datum alone. This engine hard-wires a single TOAST relation.
- `va_extsize` is the **stored** (post-compression) size and `va_rawsize` the original; when they differ, the value is compressed and the low bits of `va_rawsize` encode which algorithm. There is no chunk count — PostgreSQL derives it by scanning the TOAST index.


2. ToastChunk Structure (`[include/pg/toast.h:39-43]`):

- `toast_id` (8 Bytes), `chunk_seq` (4 Bytes), `data` (`std::vector<uint8_t>` , storing up to 2,048-byte chunks).

In PostgreSQL a TOAST table is an ordinary heap relation named `pg_toast.pg_toast_<oid>` with the schema `(chunk_id oid, chunk_seq int4, chunk_data bytea)`, plus a **unique B-tree index on** `(chunk_id, chunk_seq)`. Detoasting is therefore an index scan, and the chunks are themselves MVCC row versions with their own `xmin` / `xmax`, visible in `pg_class` and vacuumed like any other table. This engine keeps chunks in a bespoke file with a linear layout and no index.

3. Sequential Scan Preservation: Sequential scans over normal columns (e.g. `SELECT price FROM items;`) never read or load TOAST pages into shared buffers, avoiding massive I/O pollution (`[src/toast.cpp:65]`).

3. Sequence Diagram: Storing and Reassembling TOAST Data



4. PostgreSQL Fidelity Check

Claim	Verdict
Oversized values move out of line to an auxiliary relation, replaced by a small pointer	Exact
Pointer is 18 bytes	Exact size, different fields (<code>varatt_external</code>)
Value is sliced into fixed-size chunks addressed by <code>(id, seq)</code>	Exact
Scans of non-TOASTed columns never touch TOAST pages	Exact — detoasting is lazy, driven by <code>pg_detoast_datum()</code>
Threshold 2048 bytes per attribute	Wrong — <code>TOAST_TUPLE_THRESHOLD</code> is 2032 , applied to the whole tuple
Chunk size 2048 bytes	Wrong — <code>TOAST_MAX_CHUNK_SIZE</code> is 1996
No compression step	Missing — PostgreSQL compresses (pglz/lz4) before externalizing

Claim	Verdict
Single hard-wired TOAST relation, no index	Simplified — PostgreSQL uses a real heap + unique index on <code>(chunk_id, chunk_seq)</code>
No storage strategies	Missing — <code>PLAIN</code> / <code>EXTENDED</code> / <code>EXTERNAL</code> / <code>MAIN</code>
No 1 GB limit modelling	Missing — a single TOASTed varlena tops out at 1 GB in PostgreSQL


12. Item 11: SQL REPL CLI & End-to-End Engine Capstone

Confidence: verified

Citations: [include/pg/engine.h:1-85](#), [src/engine.cpp:1-250](#), [src/main.cpp:1-95](#), [tests/test_repl.cpp:1-125](#)

1. Engine Coordination Architecture

The `pg::Engine` class provides a unified facade coordinating transactions, memory buffers, slotted heap tables, secondary indexes, WAL logging, and TOAST auxiliary relations.



2. Invariants & Execution Rules


- Auto-Transaction Invariant:** Every standalone DML statement executed outside of an explicit `BEGIN ... COMMIT` block automatically receives an implicit transaction ID, snapshot, and immediate WAL commit flush (`[src/engine.cpp:38]`). **Exactly PostgreSQL's behaviour** — every statement runs in a transaction whether you asked for one or not. One difference: PostgreSQL assigns an XID *lazily*, only when a statement first writes, so read-only statements consume no transaction id at all.
- Tabular ASCII Formatting:** Query outputs render PostgreSQL-standard tabular ASCII grids displaying `item_id`, `price`, `xmin`, `xmax`, and physical `CTID`. `xmin`, `xmax` and `ctid` are genuine PostgreSQL

system columns — `SELECT ctid, xmin, xmax, * FROM items;` works against a real server — but PostgreSQL hides them unless named explicitly, whereas this engine always shows them.

3. Physical Diagnostic Commands:

- `DUMP PAGE <id>;` : Outputs exact byte offsets, `pd_lower` , `pd_upper` , line pointer flag states, and hex dumps.
- `STATUS;` : Outputs active buffer pool residency, cache hit rates, WAL flushed LSN, and transaction horizons.

3. Sequence Diagram: End-to-End Insert Pipeline



4. PostgreSQL Fidelity Check

Claim	Verdict
Every statement runs inside a transaction, implicit if not declared	Exact (PostgreSQL assigns the XID lazily on first write)
<code>xmin</code> , <code>xmax</code> , <code>ctid</code> are queryable system columns	Exact
A single coordinator object owning heap, index, WAL, buffers, CLOG, TOAST	Architecturally divergent — PostgreSQL is multi-process, one backend per connection, coordinating through shared memory

The whole upper half of PostgreSQL is out of scope here

This engine's "Lexer / Command Dispatcher" recognises a fixed statement vocabulary against a single hard-coded `items(item_id int, price int)` schema. PostgreSQL between the wire and the storage layer has:

- **A real parser** (`gram.y`) producing a parse tree, then analysis/rewrite (view expansion, rules, row-level security).
- **A cost-based planner** — join order search, `Seq Scan` vs `Index Scan` vs `Bitmap Heap Scan` , nested loop / hash / merge joins, all costed against `pg_statistic` histograms and MCV lists collected by `ANALYZE` .

- **An executor** built from pluggable plan nodes, with parallel workers and JIT compilation of expressions.
- **A system catalog** — `pg_class`, `pg_attribute`, `pg_index`, ... which are themselves ordinary heap tables read through the same buffer manager, cached in the relcache/syscache.
- **The type system** — `pg_type`, operator classes, collations, extensions.

None of that exists here, which is the point: this project is the storage engine underneath, not the database.

13. Item 12: Buffer Pool Single Gateway Invariant

Confidence: verified


Citations: [include/pg/heap.h:65-75](#), [src/heap.cpp:25-65](#), [tests/test_buffer_integration.cpp:1-125](#)

1. The Single Gateway Principle

In a production database engine, **no subsystem is ever permitted to bypass shared buffers**. Item 12 eliminates all direct `Pager::read_page()` and `Pager::write_page()` calls from `HeapFile` and `Engine`.


PostgreSQL holds the same rule for ordinary relation access, with named exceptions worth knowing:

- **Temporary tables** use per-backend *local* buffers (`localbuf.c`), not shared buffers.
- **WAL** has its own `wal_buffers` ring and never travels through the shared buffer pool.
- **Bulk operations** (`COPY`, `CREATE INDEX`, `ALTER TABLE` rewrites) go through the buffer manager but inside a restricted `BufferAccessStrategy` ring so they cannot evict the working set.
- A handful of paths talk to the storage manager (`smgr`) directly — relation extension, `mdtruncate`, and (from v17) the streaming read interface, which still lands in shared buffers but issues the I/O differently.



2. Invariants & Pin/Unpin Lifecycle State Machine

- 1. **Mandatory Symmetrical Unpinning:** Every `fetch_page()` call **must** have an exactly matching `unpin_page()` invocation in the same call stack (`[src/heap.cpp:55]`).
- 2. **Dirty Flag Propagation:** If any tuple bytes, line pointers, or page header offsets are modified in RAM, `unpin_page(page_id, is_dirty=true)` must be passed to ensure the buffer pool schedules disk writeback.



4. PostgreSQL Fidelity Check

Claim	Verdict
All relation page access flows through the buffer manager	Exact for ordinary access
Every <code>fetch_page</code> is matched by an <code>unpin_page</code>	Exact — PostgreSQL's <code>ReadBuffer</code> / <code>ReleaseBuffer</code> discipline, with <code>ResourceOwner</code> tracking to catch leaks
Dirtiness is declared by the caller at unpin time	Divergent in mechanism — PostgreSQL calls <code>MarkBufferDirty()</code> at the moment of mutation, while holding the exclusive content lock, and that call is what pairs with WAL logging

Claim	Verdict
Pin/unpin is the only lock	Simplified — PostgreSQL separates the <i>pin</i> (page stays resident) from the <i>content lock</i> (<code>LW_SHARED</code> / <code>LW_EXCLUSIVE</code> on the bytes); the state machine above conflates them
No local buffers / no strategy rings	Missing — see §1

Implementation status (2026-08-27)

The gateway is now real. A relation owns exactly one `BufferPoolManager` and there is no path around it: `HeapFile` builds one in its constructor if none is injected, and `Vacuum` and WAL recovery both go through it. Previously `VACUUM`, `CLOG` and `TOAST` each used their own `Pager` directly, which is how the cached copy and the on-disk copy drifted apart.

`PinnedPage` replaces the fetch-copy-unpin-refetch-copy-back pattern: the pin is held for the whole read-modify-write, so nothing can interleave between the read and the write-back. Dirtiness is still declared by the caller, via `PinnedPage::mark_dirty()`.

Still outside the gateway: `CLOG` and `TOAST` keep their own pagers.

14. Item 13: CLOG Commit Status 2-Bit Bitmap Pages

Confidence: verified


Citations: [include/pg/clog.h:1-55](#), [src/clog.cpp:1-110](#), [tests/test_clog.cpp:1-125](#)

1. The Commit Log (CLOG) Architecture

PostgreSQL avoids having to revisit every tuple a transaction touched at commit time by persisting one commit *state* per transaction into dedicated 8KB **commit log** bitmap pages — the `pg_xact` directory (named `pg_clog` before PostgreSQL 10; the internal API is still `clog.c`).

Each transaction is encoded using exactly **2 bits**:

00 = IN-PROGRESS 01 = COMMITTED 10 = ABORTED 11 = SUB-COMMITTED



2. Invariants & Mathematical Mapping Formulas


1. Transaction Density Derivation:

$$\text{Transactions Per Page} = 8192 \text{ bytes} \times 4 \text{ tx/byte} = 32768 \text{ transactions}$$

2. **Exact Bit Offset Formula:** $\text{clog_page} = \lfloor \text{tx_id} / 32768 \rfloor$ $\text{byte_offset} = \lfloor (\text{tx_id} \bmod 32768) / 4 \rfloor$ $\text{bit_shift} = (\text{tx_id} \bmod 4) \times 2$
3. **Instant Visibility on Restart:** CLOG pages are flushed on commit and loaded on startup, enabling instantaneous MVCC snapshot visibility checks without requiring full WAL log replay (`[src/clog.cpp:45]`).

PostgreSQL does not fsync CLOG at commit. The durability of a commit rests entirely on the WAL commit record; the CLOG bit is set in a shared-memory SLRU buffer and written out lazily, normally at checkpoint. If the server crashes between the two, redo re-applies the commit record and re-sets the bit. Fsyncing CLOG per commit would add a second synchronous write to every transaction for no durability gain.

3. Sequence Diagram: CLOG Commit and Status Lookup



4. PostgreSQL Fidelity Check

Claim	Verdict
2 bits per transaction, packed 4 per byte	Exact
32,768 transactions per 8KB page	Exact — <code>CLOG_XACTS_PER_PAGE</code>
Status codes <code>00</code> / <code>01</code> / <code>10</code> / <code>11</code> = in-progress / committed / aborted / sub-committed	Exact — <code>TRANSACTION_STATUS_*</code> in <code>clog.h</code>
Page / byte / shift arithmetic	Exact
Read-modify-write of the 2-bit field under a page pin	Exact in shape
"CLOG pages are flushed on commit"	Wrong — durability comes from the WAL commit record; CLOG is an SLRU written at checkpoint
Name "CLOG"	Dated — the directory is <code>pg_xact</code> since PostgreSQL 10

Mechanisms PostgreSQL layers on top

- **Hint bits.** The first reader that resolves a transaction's status through CLOG writes the answer back into the tuple's `t_infomask` (`HEAP_XMIN_COMMITTED` , `HEAP_XMAX_COMMITTED` , ...), so subsequent visibility checks skip CLOG entirely. Without them, CLOG would be the hottest structure in the system — which is exactly what happens in this engine.
- **SLRU segmentation and truncation.** CLOG lives in 256 KB segment files (32 pages each) under a generic simple-LRU cache. VACUUM advances `datfrozenxid` , and old segments are **deleted** — the commit log is not kept forever. Without freezing (Item 5), this engine's CLOG grows without bound.
- **SUB_COMMITTED is real.** PostgreSQL uses `0x03` for a subtransaction that has committed but whose top-level parent has not; `pg_subtrans` records the parent link. This engine reserves the code but has no subtransactions.
- **Two-phase commit.** `PREPARE TRANSACTION` state lives in `pg_twophase` , outside CLOG.

15. Item 14: On-Disk B-Tree Index with Dynamic Node Splits


Confidence: verified

Citations: [include/pg/disk_btree.h:1-95](#), [src/disk_btree.cpp:1-240](#), [tests/test_disk_btree.cpp:1-140](#)

1. Disk-Resident B-Tree Architecture

`DiskBTree` stores hierarchical B-Tree nodes on dedicated **8,192-byte disk pages**. Nodes split dynamically when overflowing, promoting median keys up to parent internal nodes.

This is a textbook B+Tree. PostgreSQL's `nbtree` is a **Lehman & Yao B-link tree**, which differs structurally — see §4.



2. Invariants & Binary Page Formats

1. B-Tree Page Header (12 Bytes — this engine's own format) (`[include/pg/disk_btree.h:20]`):

- `is_leaf` (1 Byte): `true` for leaf pages, `false` for internal routing nodes.
- `num_keys` (2 Bytes): Number of active key entries in this node.
- `right_sibling` (4 Bytes): `page_id_t` linking to the right sibling page for fast linear range scans.
- `parent_page_id` (4 Bytes): `page_id_t` of the parent router node. **PostgreSQL has no parent pointers.** A descending search pushes the blocks it passed through onto a stack, and a split walks that stack back up; this is deliberate, because a parent pointer would have to be rewritten in every child on every split, which cannot be done atomically under concurrency.


2. Entry Binary Layouts:

- **Leaf Entry (10 Bytes)** ([include/pg/disk_btree.h:39-42]): key (4B int32) + ctid (6B CTID {page, slot}).
- **Internal Router Entry (8 Bytes)** ([include/pg/disk_btree.h:44-48]): key (4B int32) + child_page_id (4B page_id_t).

3. **Median Split Algorithm:** When a leaf exceeds capacity ($N \geq 64$, BTREE_MAX_LEAF_KEYS), the median entry at $N/2$ is promoted to the parent node, splitting the keys evenly into a newly allocated 8KB sibling page ([src/disk_btree.cpp:125]).

PostgreSQL splits on **byte occupancy against fillfactor** (default 90 for leaves), not on a fixed key count, and it does not use a plain median: `_bt_findsplitloc()` scores candidate split points, and detects **rightmost-page appends** to split 90/10 instead of 50/50 so that monotonically increasing keys (a `serial` primary key) do not leave every page half empty.

3. Sequence Diagram: Leaf Split & Key Insertion



4. PostgreSQL Fidelity Check

Claim	Verdict
B+Tree on 8KB pages; keys only in internal nodes, payload in leaves	Exact in shape
Leaf pages chained for range scans	Exact in spirit — PostgreSQL chains them doubly (<code>btpo_prev</code> and <code>btpo_next</code>)

Claim	Verdict
Split allocates a sibling and pushes a separator up	Exact in shape
<code>parent_page_id</code> stored in each node	Divergent — PostgreSQL stores no parent pointers; it uses a descent stack
Split at a fixed 64-key limit, promoting the median	Divergent — PostgreSQL splits on <code>fillfactor</code> bytes, with a 90/10 rule for rightmost appends
12-byte custom node header	Divergent — PostgreSQL reuses the standard 24-byte page header and puts B-tree state in the <i>special area</i> (<code>BTPageOpaqueData</code> , 16 bytes) at the page end

What PostgreSQL's nbtree adds

- **Metapage at block 0.** Holds the root block number and tree level, so the root can move without any other page changing. The "fast root" pointer lets a mostly-empty tree skip levels.
- **High keys.** Every page stores an upper bound for its own key range as item 1. Combined with the right-link, this is what makes the tree a **B-link tree**: a reader that arrives at a page mid-split notices the searched key exceeds the high key and simply *moves right*, so searches never block on splits.
- **Page deletion is two-stage.** Emptied pages are marked half-dead, then deleted, then held until no snapshot could still be following a link into them, then recycled via the FSM. Compare: this engine never reclaims a B-tree page.
- **Deduplication (v13+).** Equal keys are folded into posting-list tuples carrying a sorted TID array.
- **Suffix truncation.** Separator keys pushed into internal pages are truncated to the shortest distinguishing prefix, so internal pages fan out further.
- **Concurrency.** Latch coupling on descent, per-page `LWLock`s, and vacuum interlocks — all absent here, where a single engine mutex serialises everything.

16. Item 15: WAL Checkpoints & Full-Page Images (FPI)


Confidence: `verified`

Citations: `include/pg/wal.h:20-60`, `src/wal.cpp:115-180`, `tests/test_checkpoint.cpp:1-130`

1. Checkpoint & Torn-Page Protection

Checkpoints truncate the recovery horizon so that crash recovery does not need to scan WAL logs from the beginning of time.

Full-Page Images (FPI) prevent catastrophic corruption caused by **torn pages** (when the OS crashes in the middle of writing an 8KB page, leaving a half-written 4KB sector on disk).



2. Invariants & Mechanics

1. **Dirty Buffer Coalescing:** `log_checkpoint()` invokes `BufferPoolManager::flush_all()`, writing all modified RAM frames to disk before appending the `CHECKPOINT` record (`src/wal.cpp:135`).

The REDO point is not the checkpoint record's LSN. PostgreSQL records the current insert position as the *REDO pointer* **before** it starts flushing buffers, flushes over a long window (`checkpoint_completion_target`, default 0.9 of the interval, to avoid an I/O spike), then writes the checkpoint record at the end and stores the REDO pointer in `pg_control`. Recovery restarts from that


earlier REDO pointer, not from the checkpoint record — because pages dirtied during the flush window may have been written before their WAL was durable. Collapsing the two, as the diagram below does, is safe here only because `flush_all()` is synchronous and nothing else is running.

2. **First-Write FPI Invariant:** The first transaction to modify any page following a checkpoint must write a full 8,192-byte snapshot of the page into WAL (`RecordType::FPI = 6`). Subsequent writes to that page only log compact delta records until the next checkpoint (`[src/wal.cpp:150]`).

Exactly PostgreSQL's rule, controlled by the `full_page_writes` GUC (on by default; safe to disable only on storage with atomic 8KB writes, e.g. ZFS). Two differences in mechanism: PostgreSQL attaches the image to the *same* WAL record that performs the modification, as a backup block inside `XLogRecordBlockHeader`, rather than emitting a separate FPI record; and it strips the page's free space ("hole" between `pd_lower` and `pd_upper`) before logging, optionally compressing the remainder (`wal_compression`). This is why WAL volume spikes immediately after every checkpoint and then decays — the standard argument for spacing checkpoints further apart.

3. **Torn-Page Healing:** During recovery, replaying an FPI record unconditionally restores the physical 8KB page byte-for-byte, healing torn sectors before delta REDO logs are applied.

3. Sequence Diagram: Checkpoint & Torn Page Healing



4. PostgreSQL Fidelity Check

Claim	Verdict
Checkpoints bound recovery time by flushing dirty buffers and recording a restart position	Exact
First write to a page after a checkpoint carries a full-page image	Exact — <code>full_page_writes</code>
FPI replay overwrites the page unconditionally, healing torn writes	Exact
FPIs are replayed before subsequent delta records for that page	Exact
Recovery restarts at the checkpoint record's LSN	Divergent — PostgreSQL restarts at the REDO pointer , captured before the flush window
Checkpoint = synchronous <code>flush_all()</code>	Simplified — PostgreSQL spreads writes across <code>checkpoint_completion_target</code> via the checkpoint process
Separate <code>FPI</code> record type	Divergent — PostgreSQL attaches backup blocks to the modifying record itself
Full 8192 bytes logged	Divergent — PostgreSQL removes the free-space hole and can compress (<code>wal_compression</code>)

Not modelled

- `pg_control`, the 8 KB file holding the last checkpoint location, the database state (`in production` / `in archive recovery` / ...), and the `BLCKSZ` /version fingerprint. Recovery begins by reading it.
- **Checkpoint triggers:** `checkpoint_timeout`, `max_wal_size`, explicit `CHECKPOINT`, shutdown, and the distinct **restartpoint** used on a standby.
- **LSN-skip during replay.** PostgreSQL compares `page.pd_lsn` against the record's LSN and skips records already reflected on the page. This engine replays unconditionally.
- **Recovery targets** — PITR, `recovery_target_time`, timelines.

Implementation status (2026-08-27)

Full-page images are now actually written. `log_fpi` existed before but had no caller outside a test, so the torn-page protection described here was dead code. `HeapFile::maybe_log_fpi` now emits one on the first modification of a page after each checkpoint, guarded by `WALManager::needs_fpi`.

`log_checkpoint` also orders its work correctly now: flush every dirty frame, `fsync` the relation, and only then append the checkpoint record — so a recovery starting at that record can trust every page before it.

Still simplified: the redo point and the checkpoint record LSN are the same here (safe only because the flush is synchronous), and the whole 8KB is logged rather than removing the free-space hole.

17. Item 16: ARIES 3-Phase Crash Recovery & UNDO Pass (cpp-postgres only)

⚠ Fidelity warning — read this first

PostgreSQL does not do this. PostgreSQL's crash recovery is **redo-only**. There is no UNDO pass, there are no compensation log records, and no `ATT`-driven rollback of uncommitted work. `cpp-postgres` implements textbook ARIES here as a teaching exercise; this chapter documents *this engine*, and §4 explains what PostgreSQL does instead and why.


The claim in earlier revisions of this chapter that "PostgreSQL employs the ARIES model" was wrong.

Confidence: `verified`

Citations: [include/pg/wal.h:70-95](#), [src/wal.cpp:160-240](#), [tests/test_undo.cpp:1-135](#)

1. The 3-Phase ARIES Recovery Algorithm


`cpp-postgres` employs the **ARIES (Algorithms for Recovery and Isolation Exploiting Semantics)** model to restore consistency after an unexpected system crash or power outage:



2. Invariants & Compensation Log Records (CLRs)


- Repeating History:** REDO is executed unconditionally for all transactions (committed and uncommitted) before UNDO runs. This guarantees that uncommitted mutations, page splits, and allocations are brought into physical memory before being rolled back ([src/wal.cpp:190]). *PostgreSQL shares this half — it also repeats history for winners and losers alike — and then simply stops.*
- Compensation Log Record (CLR):** When an operation is rolled back during the UNDO pass, a CLR (type = 7) is written to WAL with an undo_next_lsn pointer. CLRs are **never undone**, preventing infinite recovery loops if the system crashes during recovery itself ([src/wal.cpp:215]).
- CLOG Status Finalization:** Once an uncommitted transaction is rolled back, its status is finalized to ABORTED (0x02) in CLOG.

3. Sequence Diagram: ARIES Crash Recovery Walkthrough



4. PostgreSQL Fidelity Check — why PostgreSQL has no UNDO

4.1 What PostgreSQL actually does on restart



There is no Analysis pass building an Active Transaction Table for rollback purposes, no backward scan, and no CLR record type in `xlogrecord.h`.

4.2 Why it can get away with it

Because **MVCC already provides rollback**. An aborted transaction's row versions were never visible to anyone: `HeapTupleSatisfiesMVCC()` asks CLOG about `xmin`, gets `ABORTED`, and treats the tuple as though it had never existed. Undoing the physical write buys nothing — the only cost is the dead space, which VACUUM reclaims on its own schedule (Item 5).

This makes `ROLLBACK` an $O(1)$ operation in PostgreSQL, at runtime *and* at recovery: set two bits in `pg_xact`, done. A `ROLLBACK` of a transaction that updated ten million rows costs the same as one that updated none.

4.3 What it pays for that

The trade is the one this whole engine is built around: dead tuples accumulate in the heap and must be vacuumed, and the indexes bloat alongside them. An undo-based engine (Oracle, MySQL/InnoDB) instead keeps the heap clean and pays on rollback and on long-running-read reconstruction. PostgreSQL's `zheap` project prototyped exactly the ARIES-style undo design in this chapter — with an undo log, discard worker, and in-place updates — and it was never merged into core.

4.4 Verdict table

Claim	Verdict
WAL, LSNs, checkpoints, full-page images, "repeating history" in redo	Exact — PostgreSQL is genuinely ARIES-influenced here
Redo applies to committed <i>and</i> uncommitted transactions	Exact
PostgreSQL "employs the ARIES model"	Wrong — PostgreSQL implements the redo half only
Phase 1 Analysis building an ATT of losers	Not in PostgreSQL — no such pass exists
Phase 3 UNDO, backward scan, tuple restoration, <code>xmax = 0</code> rewrite	Not in PostgreSQL — aborted work is left in place and reclaimed by VACUUM
Compensation Log Records with <code>undo_next_lsn</code>	Not in PostgreSQL — no CLR record type exists
Aborted transactions end up <code>ABORTED</code> in the commit log	Exact in outcome , reached without any rollback work
Recovery is idempotent under a crash-during-recovery	Exact in outcome — PostgreSQL gets this from LSN comparison against <code>pd_lsn</code> , not from CLRs

4.5 Also not modelled

- `pg_control` **state machine** and the `DB_SHUTDOWNED` vs `DB_IN_PRODUCTION` distinction that decides whether recovery runs at all.
- `pd_lsn` **skip test**. PostgreSQL skips a redo record when `record.lsn <= page.pd_lsn`, which is what makes replay idempotent.
- **Continuous recovery**. The same code path drives archive recovery, PITR, and hot standby — recovery in PostgreSQL is not a startup-only special case.


18. Item 17: TOAST Heap + Index Full Integration

Confidence: `verified`

Citations: `include/pg/tuple.h:40-55`, `src/engine.cpp:180-245`, `tests/test_toast_integration.cpp:1-135`

1. Transparent Inline vs Out-of-Line Routing

Item 17 integrates the TOAST subsystem transparently into the main SQL execution pipeline, B-Tree index scans, and cross-file durability across restarts.



2. Invariants & Infomask Flags

- `HEAP_HASEXTERNAL` : Stamped on `TupleHeader::infomask`. Signals to sequential scans and index fetchers that the tuple's trailing bytes contain an 18-byte `ToastPointer` rather than inline text (`[include/pg/tuple.h:44]`). This engine uses `0x2000`; **PostgreSQL's** `HEAP_HASEXTERNAL` is `0x0004`, and


0x2000 is `HEAP_UPDATED` (see Item 3). In PostgreSQL the flag is a fast-path hint only — the authoritative marker is the varlena datum's own 1-byte header (`VARATT_IS_EXTERNAL`), which is what `pg_detoast_datum()` actually tests.

2. **Zero Index Decoupling Impact:** Because the secondary index stores (`{item_id} \rightarrow {CTID}`), the index node size is strictly 10 bytes regardless of whether the indexed row contains a 10-byte string or a 100-megabyte document.

This holds because the index key here is the small `item_id` . It is **not** a general PostgreSQL property: if you index the *toastable* column itself, the B-tree stores the value inline in the index tuple, detoasted (though it may stay compressed). PostgreSQL caps a B-tree entry at roughly a third of a page and raises `index row size ... exceeds btree version 4 maximum 2704` past that — the reason large-text indexing normally goes through an expression index on a hash, or a GIN full-text index.

3. **Cross-Relation Durability:** On engine restart, the main heap relation (`*_heap.db`) and auxiliary TOAST relation (`*_toast.db`) are reopened simultaneously, ensuring data consistency.

3. Sequence Diagram: Transparent TOAST Query Execution



4. PostgreSQL Fidelity Check

Claim	Verdict
TOAST is transparent to the executor; detoast happens on attribute access	Exact
A flag on the tuple marks the presence of external attributes	Exact in shape

Claim	Verdict
Index entries are unaffected when the indexed column is not the toasted one	Exact
<code>HEAP_HASEXTERNAL = 0x2000</code>	Wrong for PostgreSQL — <code>0x0004</code> ; <code>0x2000</code> is <code>HEAP_UPDATED</code>
The flag is what drives detoasting	Simplified — PostgreSQL dispatches on the varlena header (<code>VARATT_IS_EXTERNAL</code>) per datum
"Index is 100% unaffected by TOAST"	True for this schema only — indexing a toastable column stores the value in the index tuple, subject to the ~2704-byte B-tree limit
Heap and TOAST relations opened together at startup	Exact in spirit — PostgreSQL links them via <code>pg_class.reltoastrelid</code>

Not modelled

- **TOAST relations are real tables.** Chunks carry their own MVCC headers, are vacuumed, and appear in `pg_class` / `pg_toast`. An `UPDATE` that does not touch the toasted column **reuses the existing chunks** — PostgreSQL copies the pointer into the new row version rather than rewriting the value, which is a large part of why wide-column updates are cheap.
- **Slice access.** `substr()` on an `EXTERNAL`-storage column fetches only the chunks it needs, via the `(chunk_id, chunk_seq)` index. This engine always reassembles the whole value.
- **Compression state.** PostgreSQL tracks compressed vs raw size in the pointer itself (Item 10).

19. Item 18: Embedded HTTP/REST API Server (Approach A)

Confidence: verified


Citations: [include/pg/http_server.h:1-45](#), [src/http_server.cpp:1-210](#), [tests/test_http_server.cpp:1-125](#)

1. 2-Tier Architecture & Web Integration

Item 18 implements **Approach A: Direct Embedded HTTP/REST Server**. The database binary embeds a lightweight, non-blocking Winsock2 HTTP 1.1 server listening on port 8080, allowing web applications (e.g. React, Vue, cURL) to query the database directly via JSON.

This has no PostgreSQL counterpart — PostgreSQL speaks only the frontend/backend protocol on 5432 (Item 19), and HTTP access is provided by external projects (PostgREST, pgHttp, Supabase's API layer) sitting in front of it. Item 18 is a convenience for the desktop demo, not a modelled PostgreSQL subsystem.

Note that `Access-Control-Allow-Origin: *` on an endpoint that executes arbitrary SQL with no authentication is only acceptable because this daemon is intended to bind loopback on a developer machine. It must not be exposed on a routable interface.



2. Invariants & REST Endpoints

1. **Full CORS Support**: Automatically handles `OPTIONS` requests and emits headers allowing cross-origin requests from any local frontend port:

- `Access-Control-Allow-Origin: *`
- `Access-Control-Allow-Methods: GET, POST, OPTIONS, PUT, DELETE`
- `Access-Control-Allow-Headers: Content-Type, Authorization`


2. **JSON Response Contract**:

- `POST /api/sql`: Executes arbitrary SQL and returns `{ "success": true, "sql": "...", "output": "..." }`.
- `GET /api/items`: Streams live table rows directly as structured JSON objects:

```
{
  "items": [
    { "item_id": 100, "price": 10, "xmin": 1, "xmax": 0, "ctid": "(0, 1)" }
  ]
}
```

3. **Thread Safety:** All incoming HTTP worker threads synchronize access to `pg::Engine` using `std::mutex engine_mutex_`.

3. Sequence Diagram: React Browser Query via HTTP



4. PostgreSQL Fidelity Check

Not applicable — no PostgreSQL analogue. See Item 19 for the subsystem that *is* protocol-accurate.

20. Item 19: PostgreSQL Wire Protocol Server (Approach C)

Confidence: verified

Citations: [include/pg/pgwire.h:1-55](#), [src/pgwire.cpp:1-320](#), [tests/test_pgwire.cpp:1-160](#)

1. Enterprise 3-Tier Integration Architecture


Item 19 implements **Approach C: Native PostgreSQL Frontend/Backend Protocol v3.0**. Listening on TCP port `5432`, `PgWireServer` transforms `cpp-postgres` into a drop-in PostgreSQL server compatible with standard PostgreSQL CLI tools, GUI clients, drivers, and ORMs.



2. Invariants & Packet Framing

- Big-Endian Integer Serialization:** All 16-bit and 32-bit fields are transmitted in Network Byte Order (`htonl` / `htons`) (`[src/pgwire.cpp:15]`).
- Binary Message Format:**
 - StartupMessage -> AuthenticationOk (R, len 8, 0) -> ParameterStatus (S) -> ReadyForQuery (Z, I).
 - A real PostgreSQL backend also sends **BackendKeyData (K , PID + secret)** between the last `ParameterStatus` and `ReadyForQuery` ; it is what a client needs in order to issue a `CancelRequest` on a second connection. Most drivers tolerate its absence, but `pg_cancel_backend` -style cancellation cannot work without it.
 - Query (Q) -> RowDescription (T) -> DataRow (D)* -> CommandComplete (C) -> ReadyForQuery (Z).
- Transaction State Synchronization:** The ReadyForQuery (Z) packet transmits I when idle and T when an active transaction block is open. PostgreSQL defines a third state, `E` — in a transaction block that has hit an error, where every statement is rejected until `ROLLBACK` . That is the state `psql` reports as `!` .

3. Sequence Diagram: PostgreSQL Wire Protocol Handshake & Query



4. PostgreSQL Fidelity Check

Claim	Verdict
Protocol v3.0, <code>SSLRequest</code> code 80877103, <code>N</code> to decline TLS	Exact
Message framing: 1-byte type tag + int32 length (length includes itself, excludes the tag)	Exact
Big-endian / network byte order throughout	Exact

Claim	Verdict
Simple Query flow Q -> T -> D* -> C -> Z	Exact
Z carries I / T transaction status	Incomplete — E (failed transaction block) also exists
Startup: R -> S * -> Z	Incomplete — a real backend sends K (BackendKeyData) before Z

Not implemented

- **Extended query protocol** (P Parse / B Bind / D Describe / E Execute / S Sync) — used by every driver that supports prepared statements or server-side parameter binding, including JDBC and psycopg by default.
- **Authentication beyond AuthenticationOk** — no scram-sha-256 , no md5 , no TLS. Any client that connects is trusted.
- **COPY sub-protocol** (G / H / d / c), N NoticeResponse, A NotificationResponse (LISTEN / NOTIFY), CancelRequest .
- **Catalog queries.** psql backslash commands (\d , \dt) issue pg_catalog queries; without those tables they will fail even though the connection succeeds.

21. Item 20: Native Desktop GUI & Unified Server Daemon


Confidence: verified

Citations: [src/gui/main_gui.cpp:1-480](#), [src/server_main.cpp:1-75](#), [CMakeLists.txt:100-150](#)

1. Unified Applications Overview

Item 20 delivers two standalone executable artifacts:

- pg_gui.exe (3.89 MB)**: A self-contained Windows desktop GUI application built with **Dear ImGui + DirectX 11 + Win32**, providing visual memory inspectors for 8KB slotted pages, clock-sweep buffer frames, on-disk B-Tree graphs, and an interactive SQL scratchpad.
- pg_server.exe** : A multi-protocol daemon running both **Approach A (HTTP REST :8080)** and **Approach C (pgwire TCP :5432)** concurrently on background threads, wired to a shared **pg::Engine** .



2. Desktop GUI Feature Breakdown

CPP-POSTGRES STORAGE ENGINE DESKTOP EXPLORER

SQL QUERY WORKSPACE

SELECT * FROM items;

[▶ Run F5] [BEGIN]

[COMMIT] [ROLLBACK]

[STATUS] [VACUUM]

LIVE TABLE VIEW

id	price	xmin	CTI
100	20	4	0,3
200	5	2	0,2

OUTPUT LOG CONSOLE

PHYSICAL STORAGE ENGINE VISUALIZERS

[8KB Slotted Page] [Shared Buffers] [Disk B-Tree]

Page 0 Header (18B) | pd_lower: 30B | pd_upper: 8120B

[Line Pointers]

Free Space

[Tuples]

Slot #	Offset	Length	Live Tuple Details
Slot 1	8168	24 B	[LIVE] id=100, price=10
Slot 2	8144	24 B	[HOT REDIRECT] -> (0, 3)
Slot 3	8120	24 B	[HEAP-ONLY TUPLE] id=100, 20

Shared Buffers Tab: Live Clock-Sweep & Pin Tracking


Disk B-Tree Tab: Key -> Candidate CTID Mapping

CLOG Tab: 2-bit Transaction Status Bitmaps

WAL Tab: Crash Recovery & Checkpoints (see Item 16)

TOAST Tab: 2KB Chunk Auxiliary Table Inspector

3. Sequence Diagram: Dual Server Daemon Operation



4. PostgreSQL Fidelity Check

Both binaries are original to this project; PostgreSQL ships `postgres` / `postmaster` plus `psql` , and its GUIs (pgAdmin, DBeaver) are separate projects speaking the wire protocol.

Two labels in the diagrams above are worth reading precisely:

- **"ARIES undo"** describes *this engine*. PostgreSQL's WAL is ARIES-influenced but redo-only — no undo pass, no CLRs (Item 16).
- **"Page 0 Header (18B)"** in the GUI mock is this engine's header. PostgreSQL's is 24 bytes (Item 2).

The `pg_server.exe` process model — one process, two listener threads, one shared `pg::Engine` behind a mutex — is also structurally unlike PostgreSQL, which forks a dedicated backend process per connection and coordinates through shared memory and lightweight locks. The single global mutex here means exactly one statement runs at a time, so none of the MVCC concurrency documented in Item 4 is actually exercised in parallel.